

# 30 Years of Engineering Multi-Agent Systems: What and Why (EMAS)

Michael Winikoff  
AAMAS 2024  
Auckland, New Zealand

These slides:  
[tinyurl.com/  
AAMAS2024Thu9am](https://tinyurl.com/AAMAS2024Thu9am)



Note: bibliography at end

# Agenda

- An overview of EMAS (Engineering Multi-Agent Systems)
  1. What
  2. Why
  3. Challenges and future directions
- Some EMAS research

These slides:  
[tinyurl.com/  
AAMAS2024Thu9am](https://tinyurl.com/AAMAS2024Thu9am)



# Agenda

- An overview of EMAS (Engineering Multi-Agent Systems)
  1. What
  2. Why
  3. Challenges and future directions
- Some EMAS research

**Disclaimer ...**

These slides:  
[tinyurl.com/  
AAMAS2024Thu9am](https://tinyurl.com/AAMAS2024Thu9am)



# Agenda

- An overview of EMAS (Engineering Multi-Agent Systems)
  1. What
  2. Why
  3. Challenges and future directions
- Some EMAS research

Disclaimer ...

Thanks:  
Lin Padgham

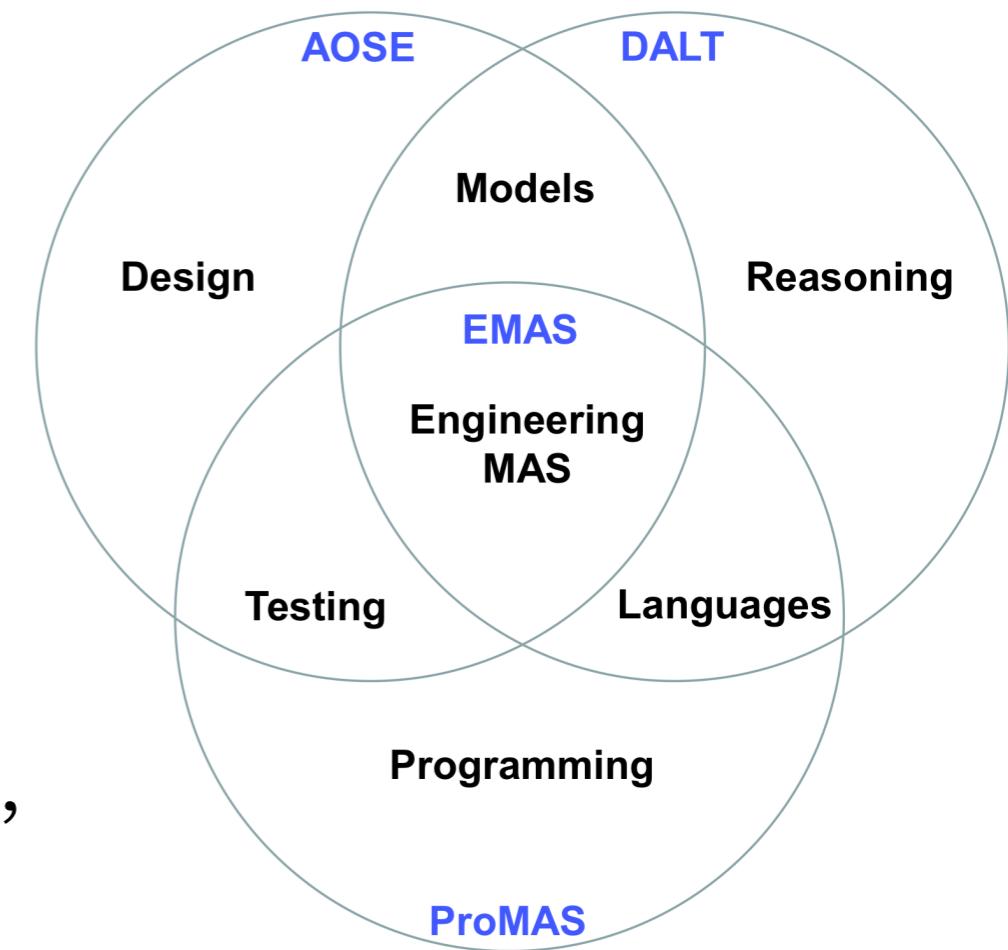


# What is EMAS?

- “Despite the substantial body of knowledge and expertise developed in the design and development of Multi-Agent Systems (MAS), the **systematic development of large-scale and open MAS** still poses many challenges” ([EMAS 2024 CFP](https://emas.in.tu-clausthal.de/2024/cfp/), emphasis added)

<https://emas.in.tu-clausthal.de/2024/cfp/>

- Research on Engineering Multi-Agent Systems is concerned with a range of topics that aim to provide software engineers with the **concepts, processes, notations, techniques, languages, and tools** to be able to effectively develop multi-agent systems.



**EMAS 2014 keynote:**  
Twenty Years of Engineering MAS  
(Hindriks, 2014)

# EMAS Work ...

- Less about algorithms/theories/theorems ...
- ... more about (human-oriented) models/processes/tools
- Importance (and challenge!) of artefacts (effort, sharing, incentives)

# Why?

- Why (multi-)agents?

- Wide range of applications involving distribution, adaptability, flexibility, robustness, autonomy, complexity

- Domains include production scheduling, energy, transport, disaster relief, manufacturing, simulation, health, UAVs ...

- Why Engineering?

- Why not just write it Java?

Benfield et al, 2006  
Belecheanu et al, 2006  
Munroe et al, 2006  
Dignum & Dignum, 2010  
Müller & Fischer, 2014  
Singh & Padgham, 2015  
Briola et al, 2023



Source: Singh & Padgham, 2015

# Benefits?

“Based on this analysis, agent-oriented methodology, architecture, and development delivered a 368% improvement on overall project productivity. (2.11 FP Actual/0.45 FP Expected) ...

In a **wide range of complex business applications**, we show that the use of **BDI** [belief-desire-intention] **technology** incorporated within an enterprise-level architecture can **improve** overall developer **productivity** by an average [of] 350%. For java coding alone, the increase in productivity was over” (Benfield et al., 2006; emphasis added)

- Paper by Agentis Software staff – company commercialising agent technology, industries include finance, insurance, logistics, energy, IT infrastructure.

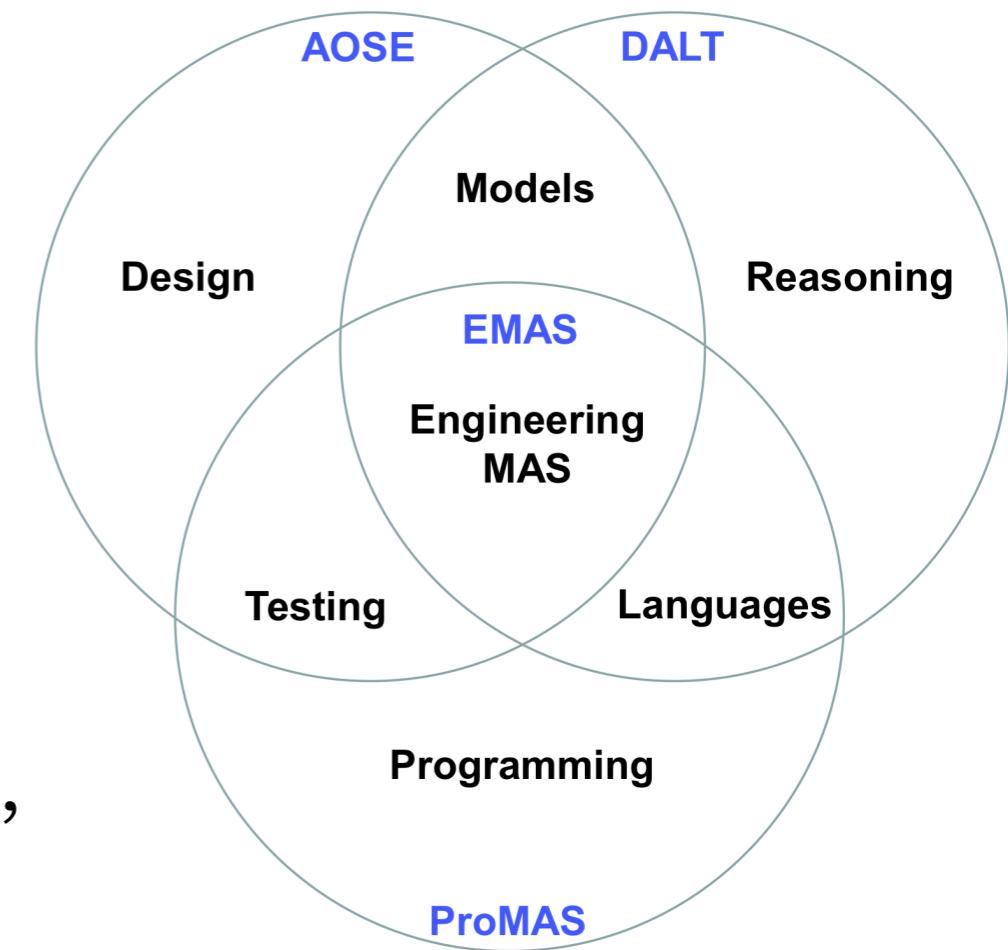
***But: unverified claims by industry ...***

# What is EMAS?

- “Despite the substantial body of knowledge and expertise developed in the design and development of Multi-Agent Systems (MAS), the **systematic development of large-scale and open MAS** still poses many challenges” ([EMAS 2024 CFP](https://emas.in.tu-clausthal.de/2024/cfp/), emphasis added)

<https://emas.in.tu-clausthal.de/2024/cfp/>

- Research on Engineering Multi-Agent Systems is concerned with a range of topics that aim to provide software engineers with the **concepts, processes, notations, techniques, languages, and tools** to be able to effectively develop multi-agent systems.

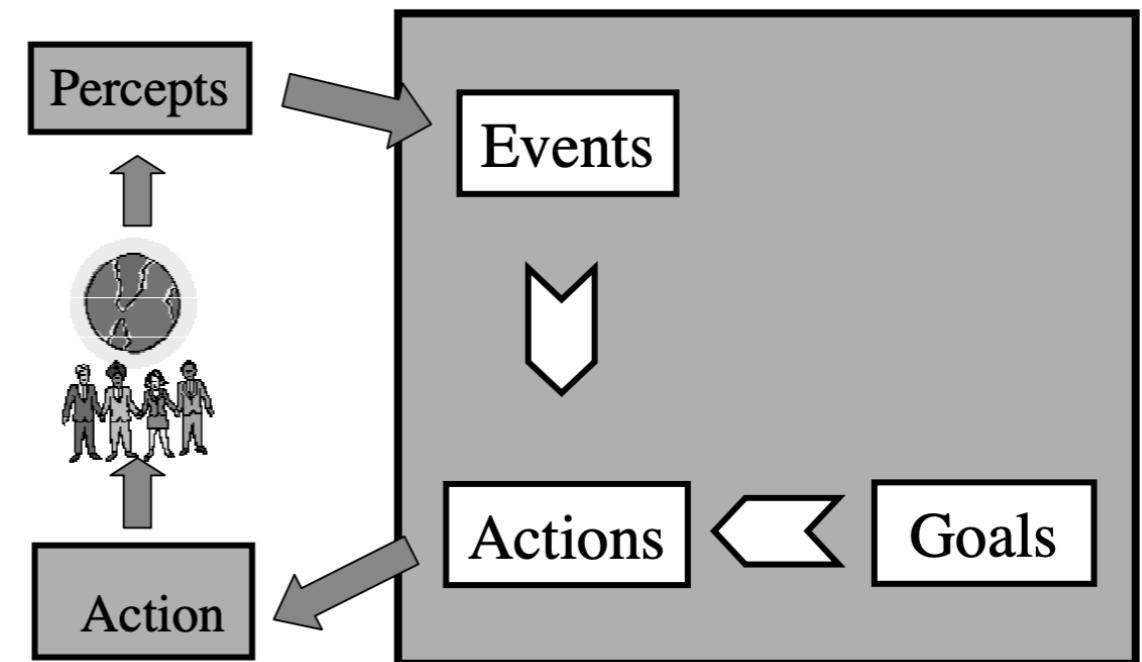


**EMAS 2014 keynote:**  
Twenty Years of Engineering MAS  
(Hindriks, 2014)

# Concepts

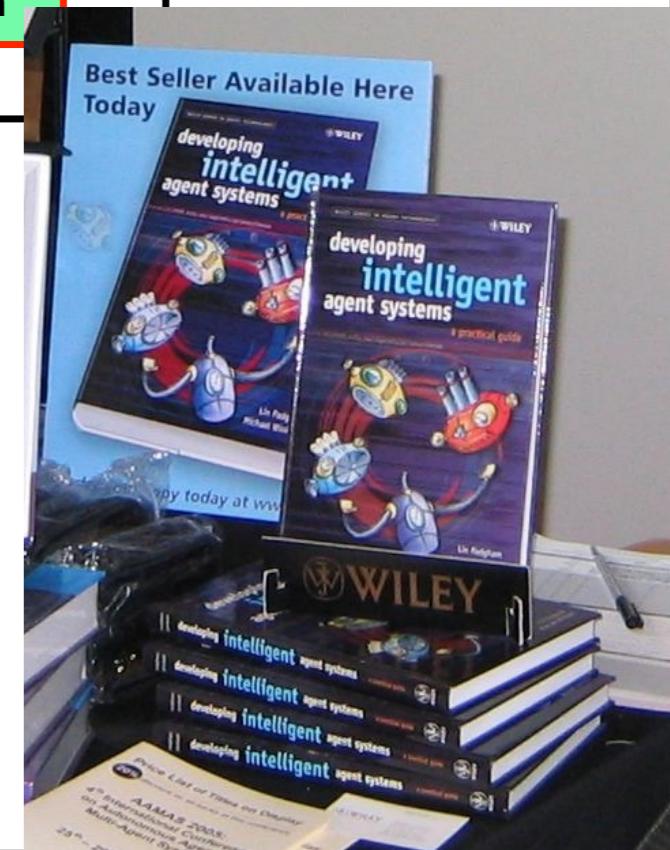
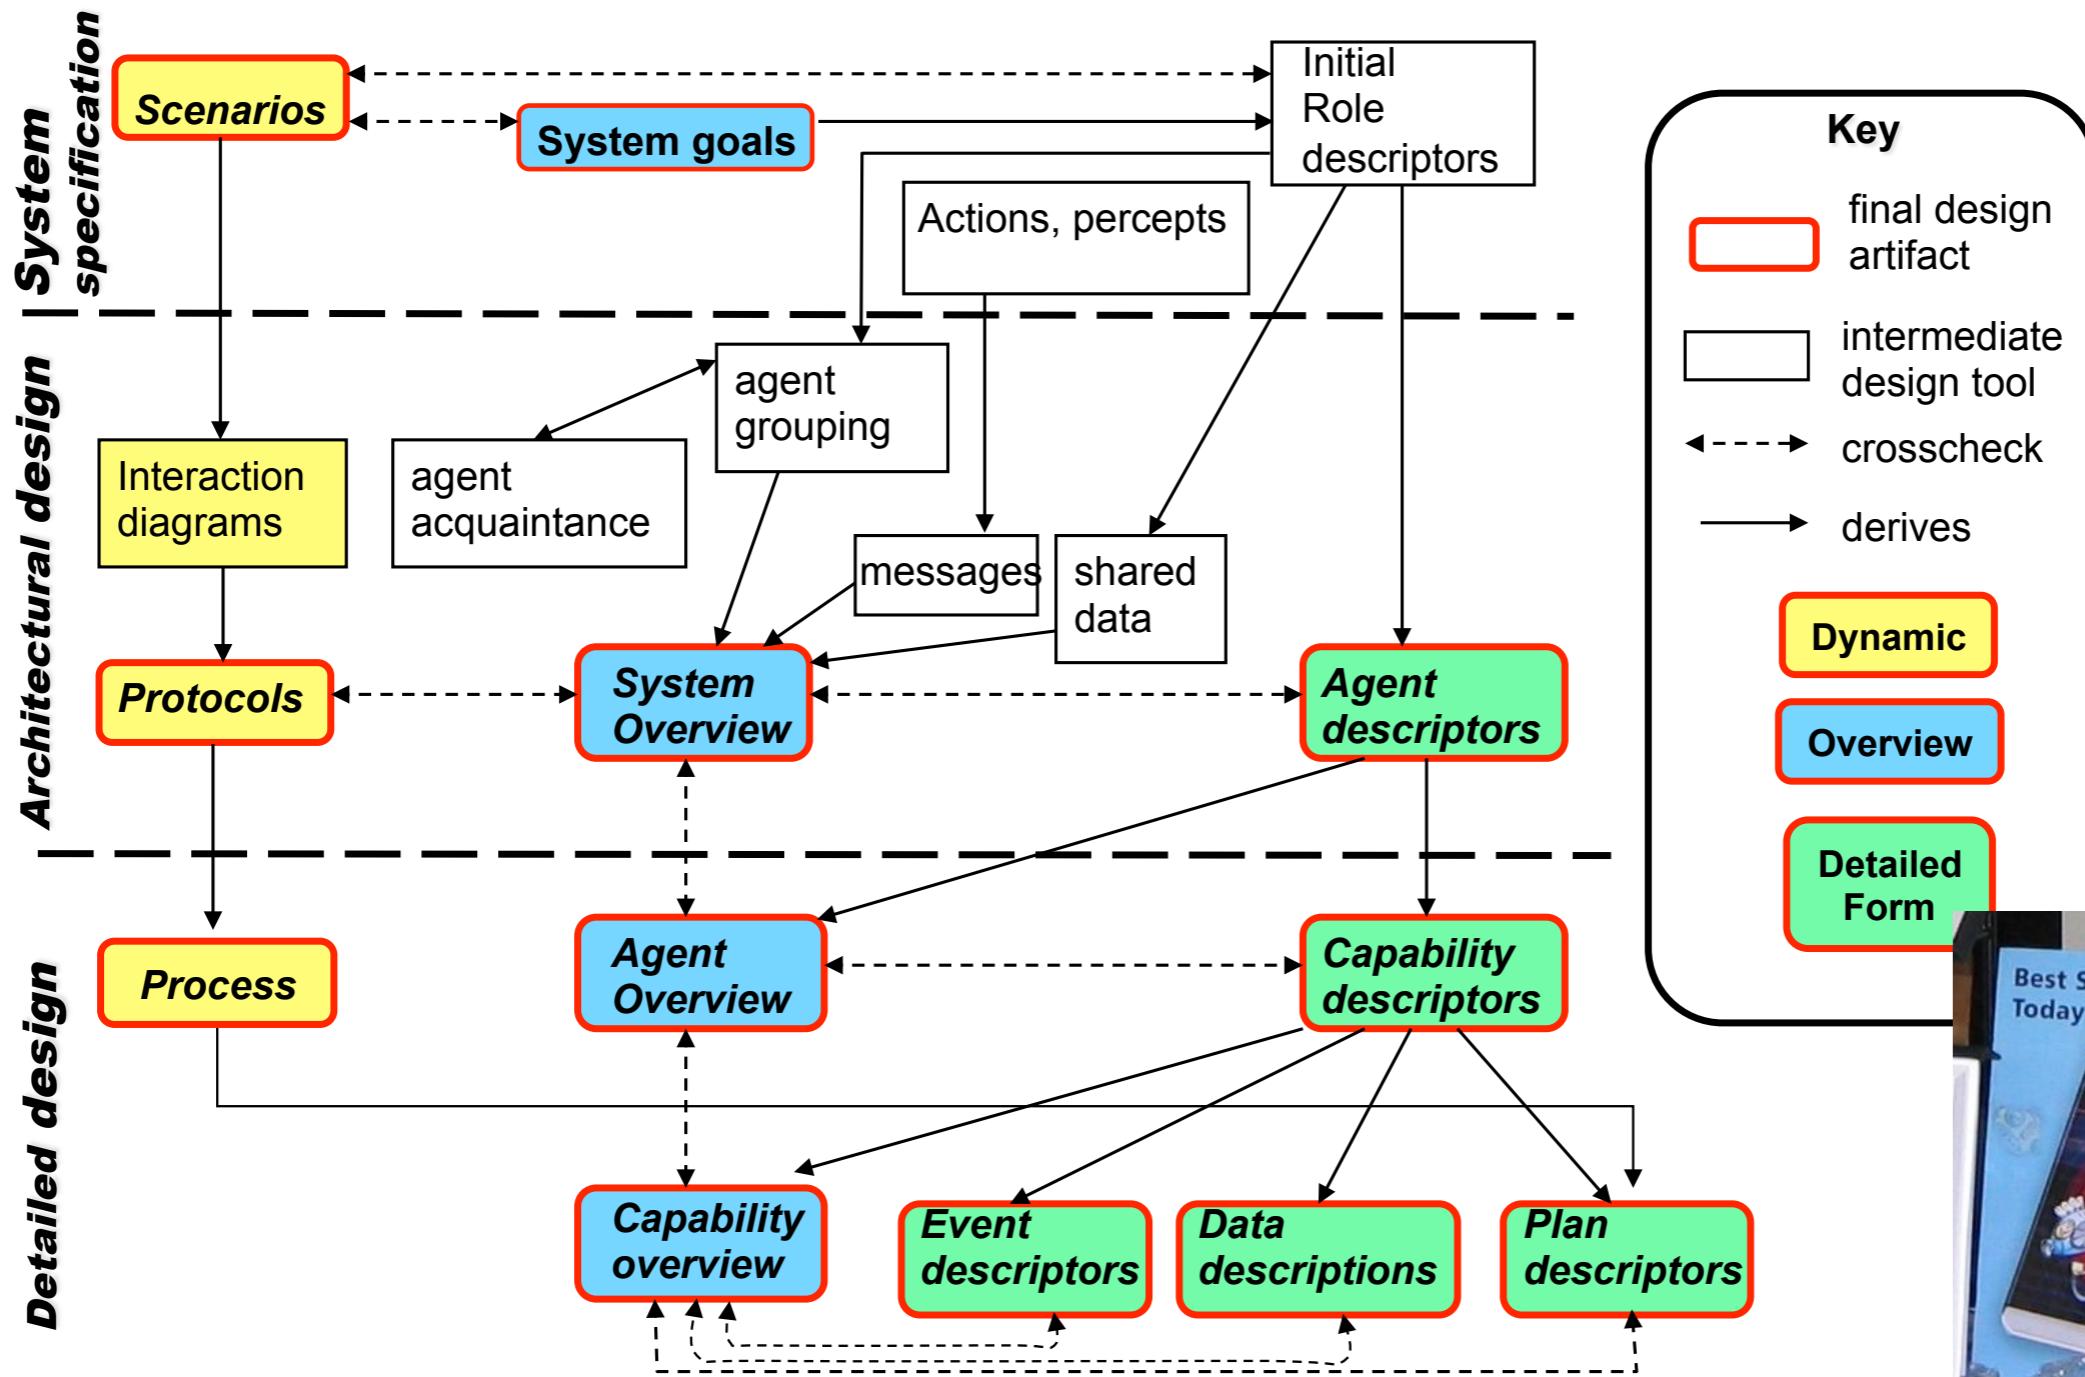
## Agents are:

- *situated* - hence **actions** and **percepts**
- *reactive* - hence **events**
- *proactive* - hence **goals** (key concept)!
- *social* - hence **messages**, **commitments**, ...



Source: Winikoff et al, 2001

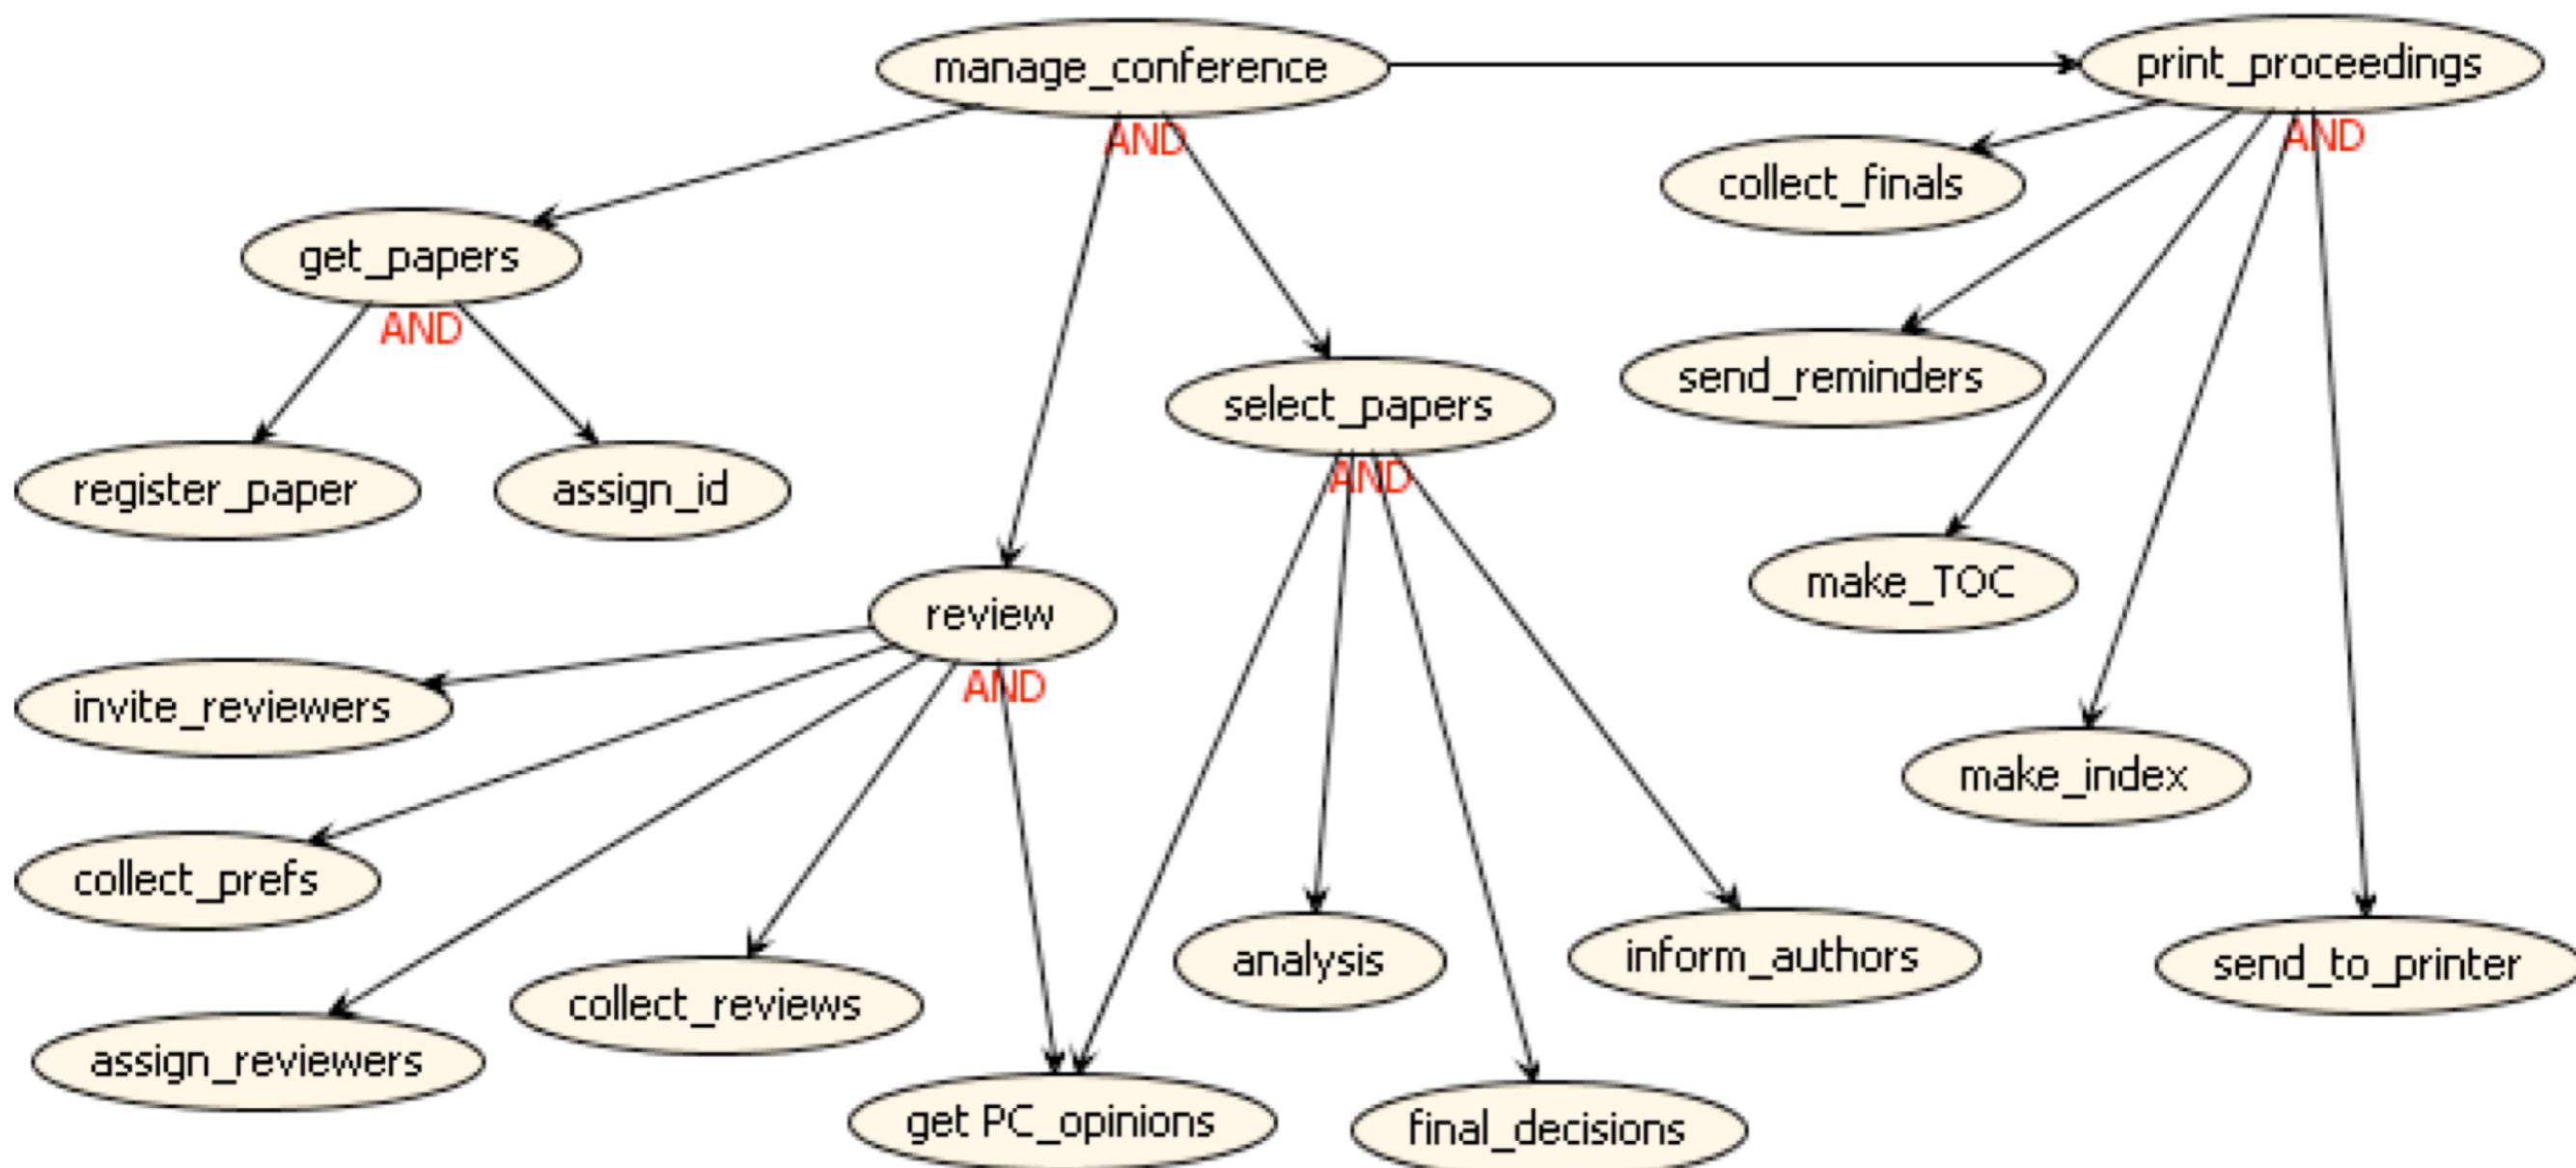
# Process



Source: Padgham & Winikoff, 2004

# Process (cont)

- ...
- Implementation
- Assurance (testing, debugging, verification)
- Maintenance
- *Note: not sequential - iterative ... (and Agile ...)*



Source: Padgham et al, 2008

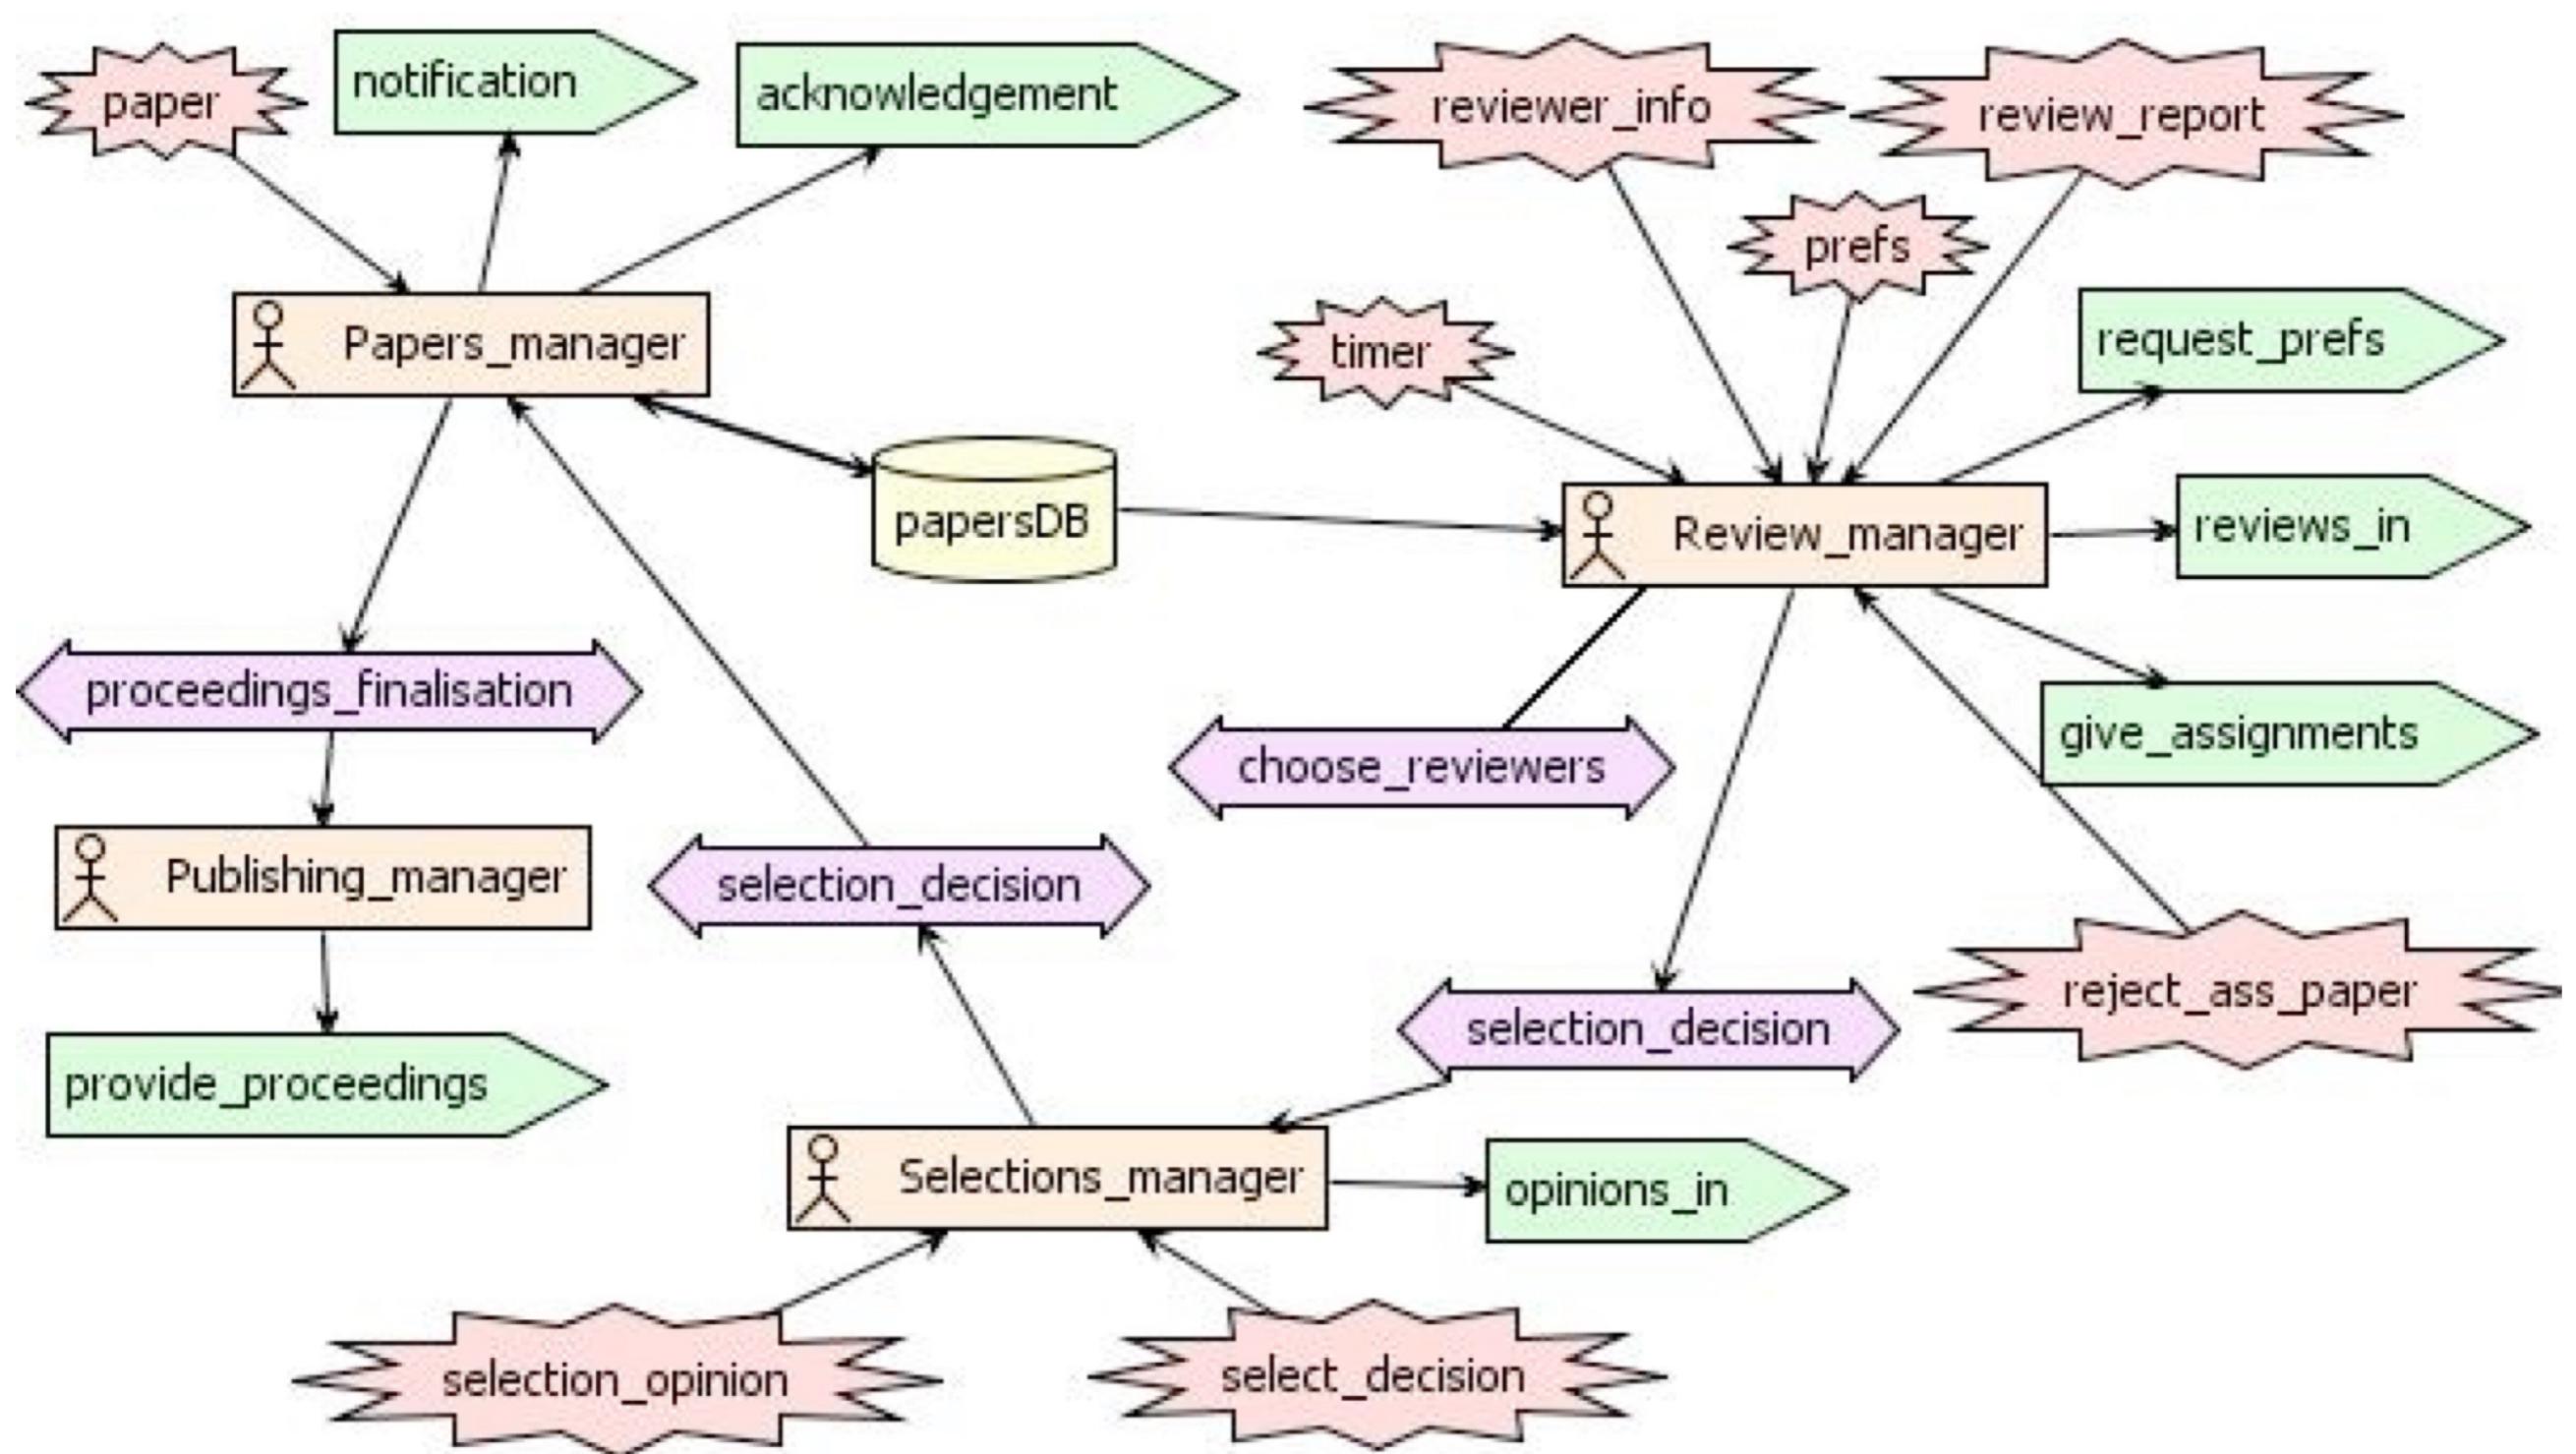
Edit Scenario - review scenario

	Type	Name	Role	Data	Description
1	G	invite_reviewers	Review_management	ReviewerDB	Invite candidates to join the review panel
2	G	collect_prefs	Assignment	ReviewerDB	Collect the preference of the reviewers
3	G	assign_reviewers	Assignment	ReviewerDB	Assign papers to reviewers based on their prefe...
4	A	give_assignments	Assignment		Send the papers to the allocated reviewers
5	P	review_report	Review_management		Receive the review from the reviewers
6	G	collect_reviews	Review_management	ReviewDB	Collect all the reviews from the reviewers

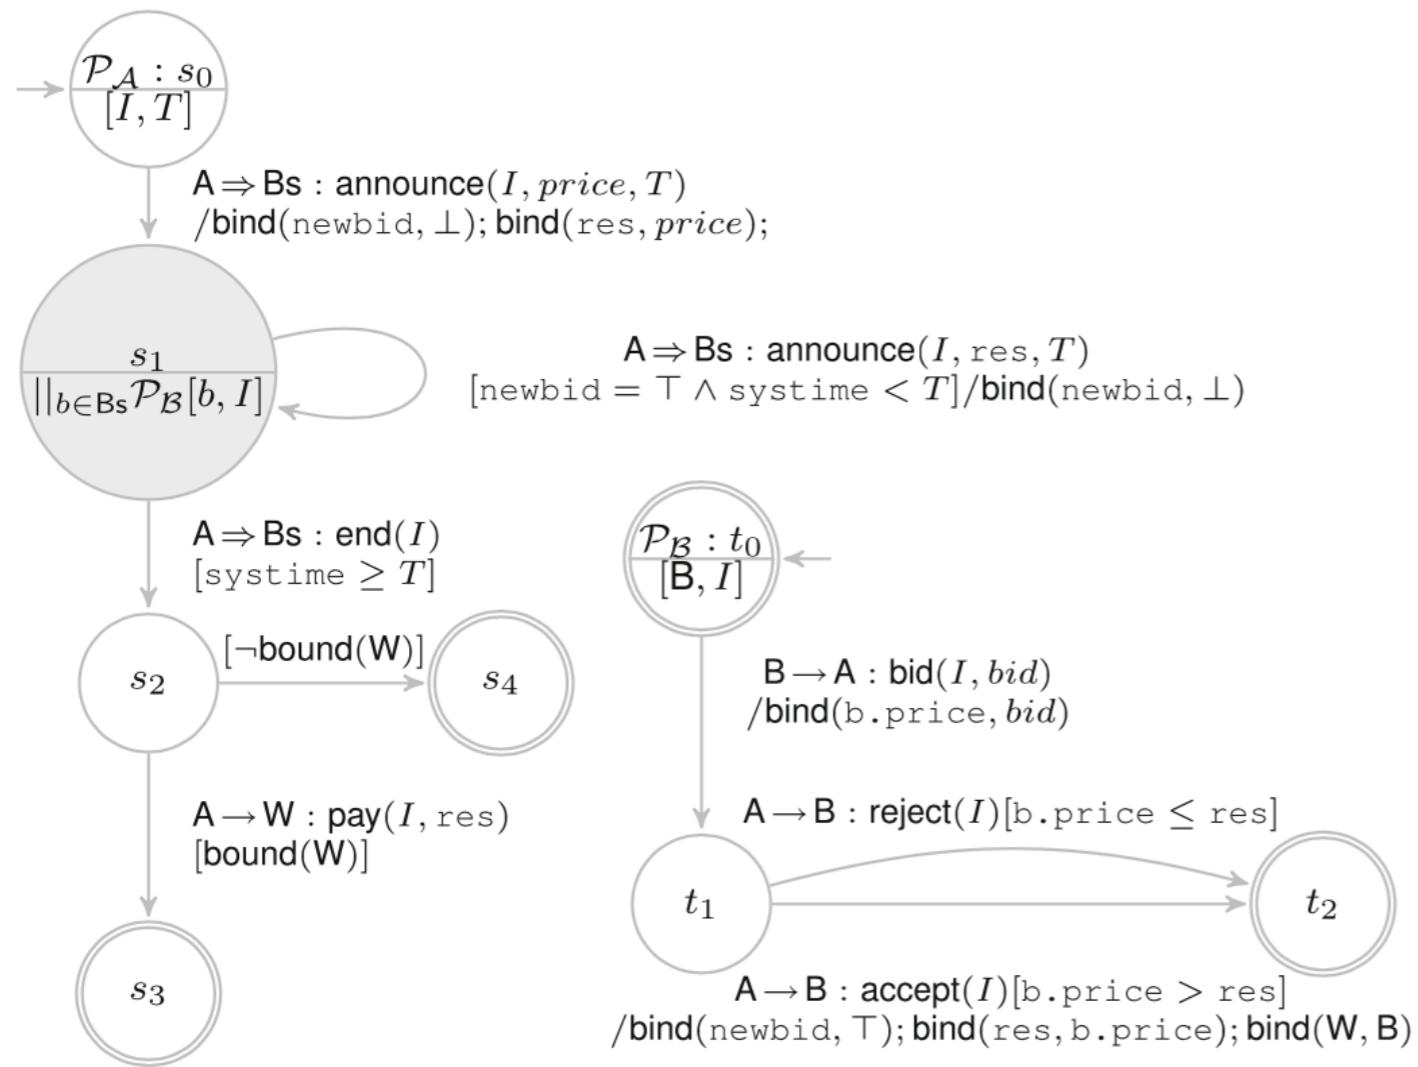
Insert Step
Edit
Remove
  
Close

A -> Action      G -> Goal      O -> Others      P -> Percept      S -> Scenario

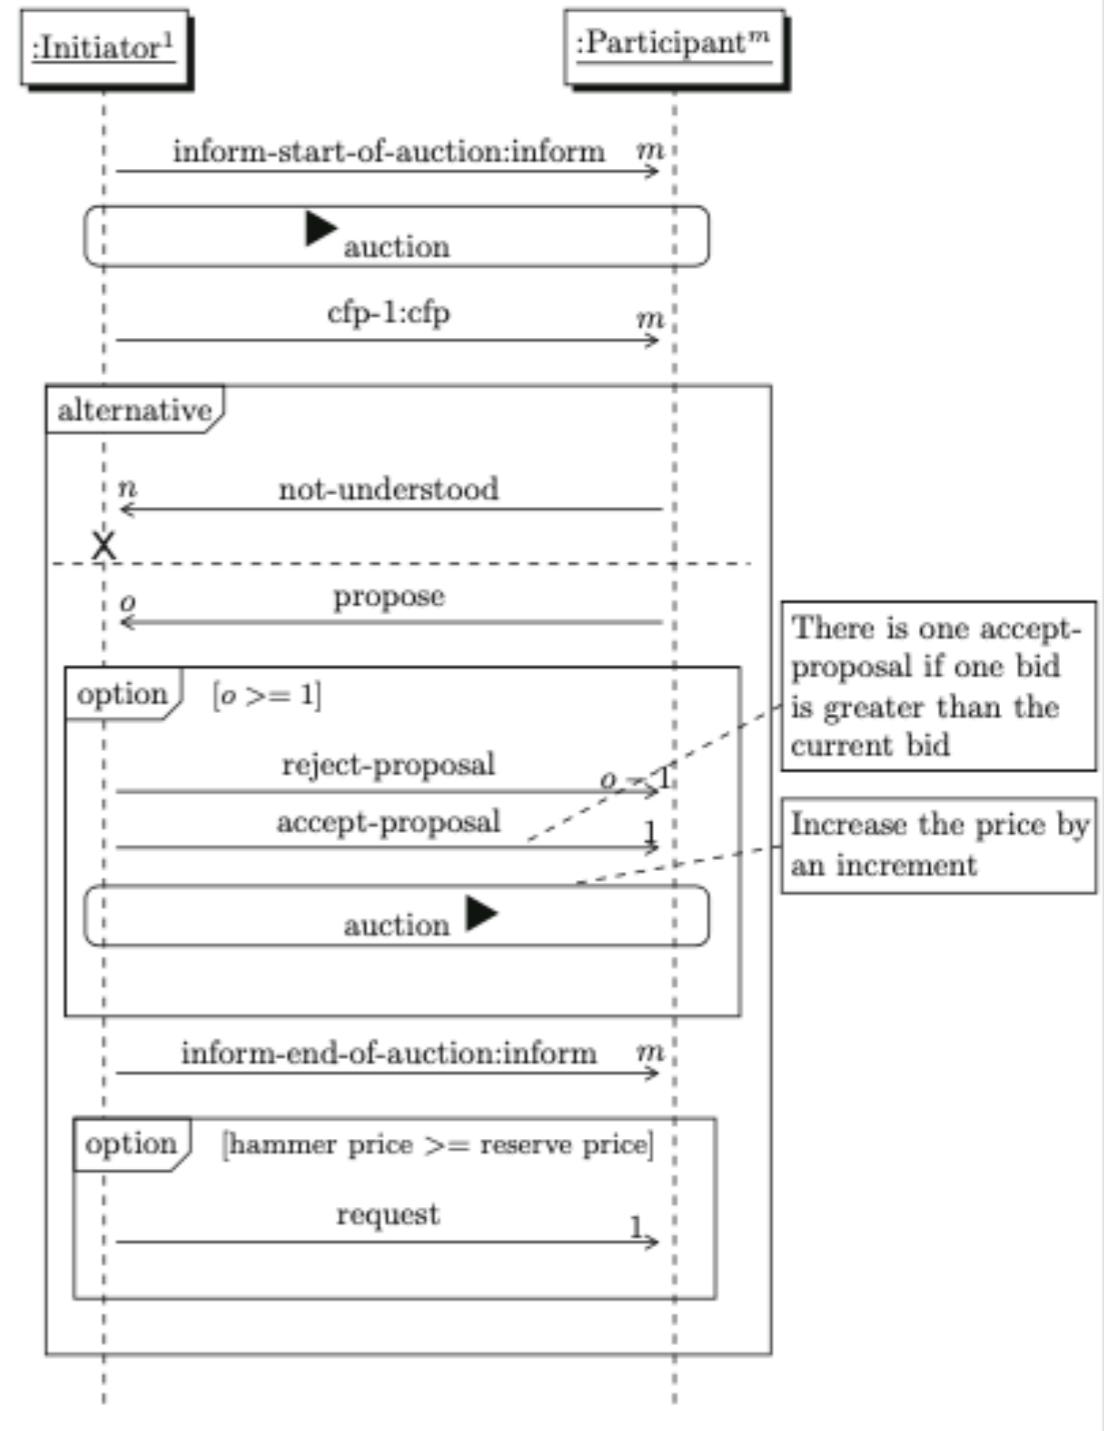
Source: Padgham et al, 2008



Source: Padgham et al, 2008



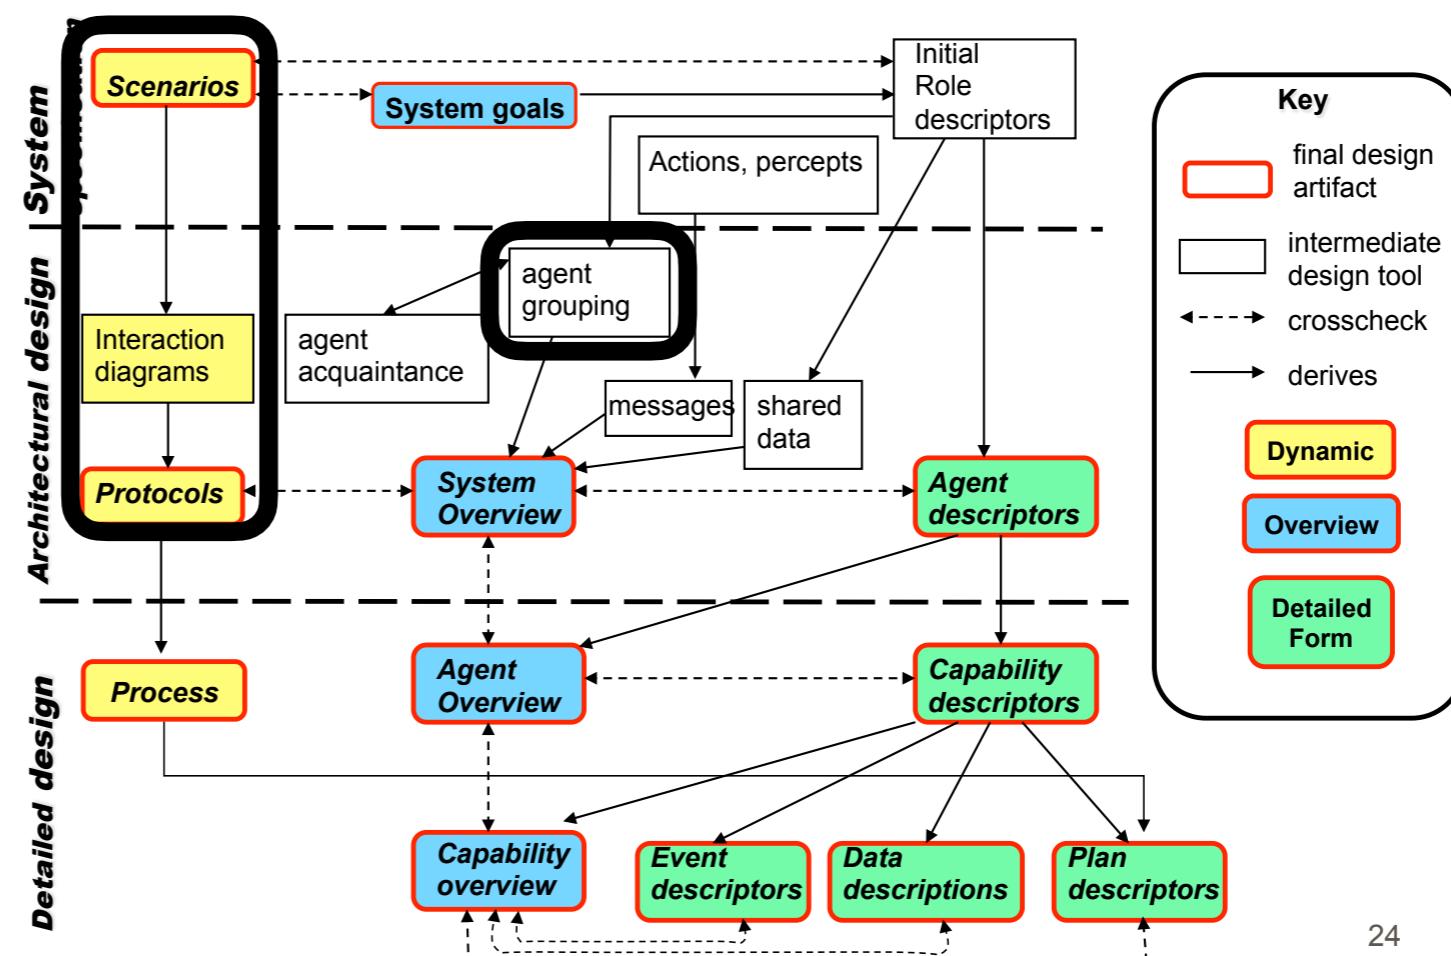
### sd FIPA English Auction Protocol



Source: Winikoff et al, 2018

# Techniques

- Example: Define agent types by grouping smaller chunks, considering coupling and cohesion
- Example: Scenarios → Interaction Diagrams → Protocols



24

# Languages

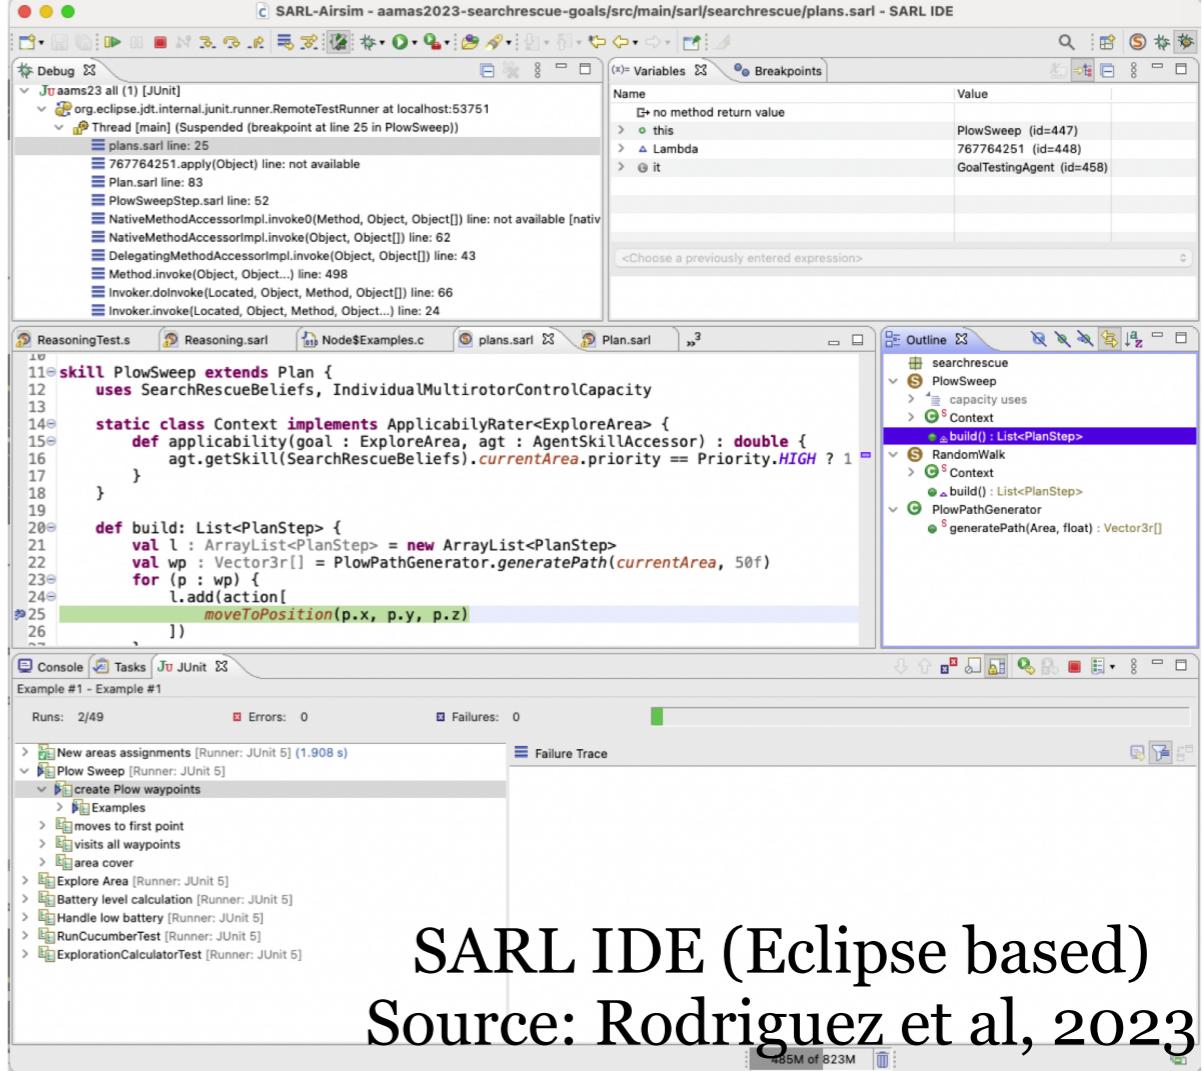
- *Could have a whole talk just on this!*
- Many languages (e.g. 2APL, 3APL, AgentSpeak, Jason, JaCaMo [Jason+CArtAgO+Moise], Astra, Brahms, GOAL, GOLOG, ConGolog, IndiGolog, Gwendolen, JACK, JADE, Jadescript, Jadex, JIAC, MetateM, SARL ...)
- Will introduce BDI (Belief-Desire-Intention) and show an example AgentSpeak program later ...

Bordini et al, 2005 & 2009  
Bordini & Dix, 2016

Bratman, 1987  
Ingrand et al, 1992

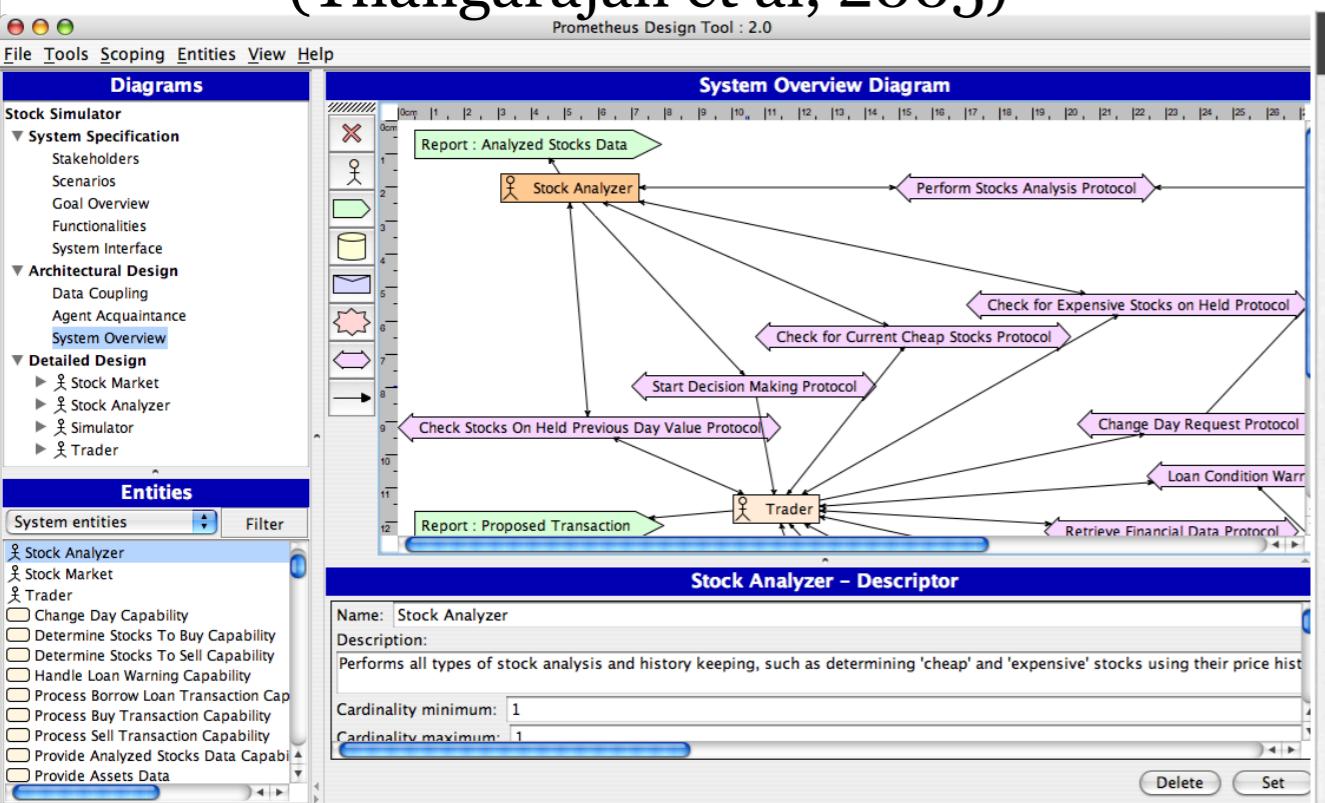
# Tools

- Support design, coding, testing, debugging, ...

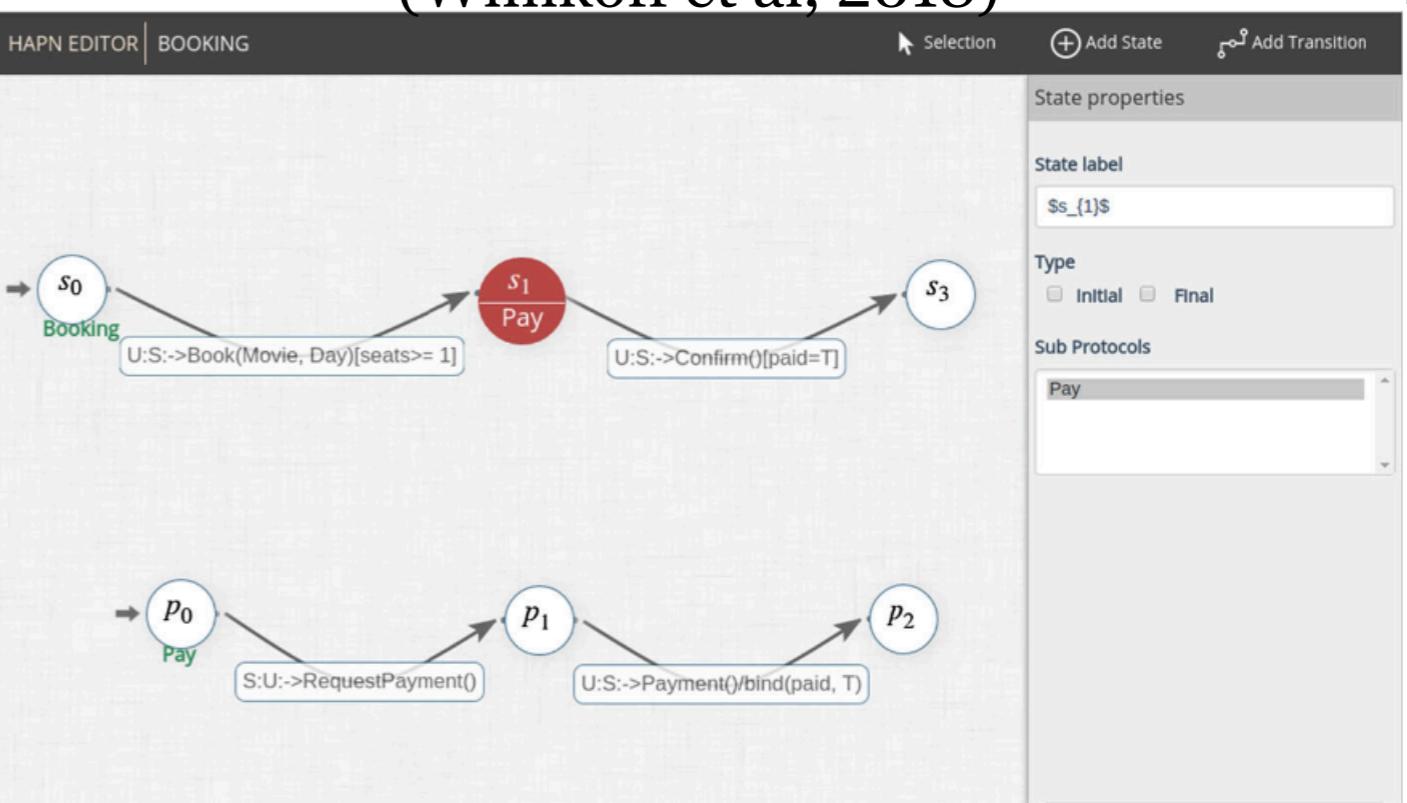


SARC IDE (Eclipse based)  
Source: Rodriguez et al, 2023

Prometheus Design Tool  
(Thangarajah et al, 2005)



HAPN Protocol Editor  
(Winikoff et al, 2018)

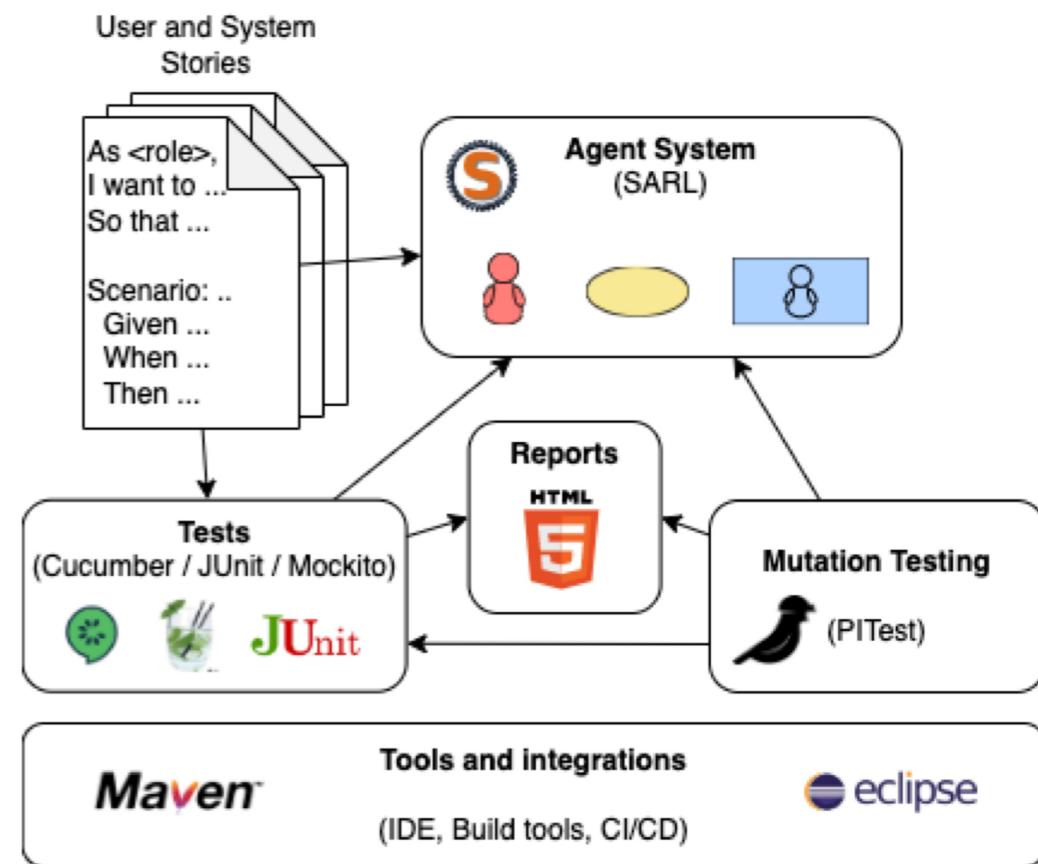


# Some EMAS Challenges

- Practicalities
- Community
- Research

# Practicalities

- **Modernise** - the world has changed! (Agile, Python, DevOps, hypermedia...)
- **Integrate** (Eclipse, etc.)
- Why? Ease-of-use (familiar tools), easier adoption, performance, robustness, scalability ...



Source: Rodriguez et al, 2023

## Rodriguez EMAS 2024 keynote

**Sebastian Rodriguez: Agile Approach for Agent Oriented Software Engineering**

(Chair: Brian Logan)



**Abstract:** The agile software development life cycle is widely used in industry today due to its highly flexible and iterative processes that facilitate rapid prototyping. In this talk, we advocate to close the gap between mainstream software engineering and agent technology by adopting and adapting well-known and accepted techniques. We introduce an agile approach to capturing requirements in agent systems via user and system stories and how to translate these requirements into goal-oriented agent models. We cover how to define test cases to verify the expected system behaviour following a Behaviour-Driven Development (BDD) approach. We leverage a range of state-of-the-art development tools, inheriting the rich set of features they provide. Finally, we discuss future directions and opportunities.

# Community

- **Link** with other relevant research and practice communities (including industry!)
- E.g. software engineering, internet/web of things, (social) AI, service, grid, autonomic ... and AAMAS beyond EMAS!
- Quantify (and market) benefits and applicability (test case, competitions ...)

# (some) Research

- **Beyond BDI:** more powerful cognition (e.g. adding values, planning, ...) - want simplicity, transparency and efficiency - difficult! (Logan, 2015)
- **What about Learning?** “Agent programming in the cognitive era” (Ricci, 2022; Bordini et al, 2020)
  - *The “Cognitive Hourglass”: Agent Abstractions in the Large Models Era*, Ricci et al., Blue Sky 2, Friday 2-4pm
- **Flexible interaction:** Commitment Machines (Yolum & Singh, 2002), BSPL (Singh, 2011; Chopra & Christie V, 2023)
- **Assurance ...**

# Assurance is important

“The verification of such complex software is **critical to its acceptance** by science mission managers.” (Havelund et al., 2000)

“... validation through extensive tests was mandatory ... However, the task proved challenging for several reasons. First, agent-based systems explore realms of **behaviour outside people's expectations and often yield surprises** ...” (Munroe et al., 2006, section 3.7.2)

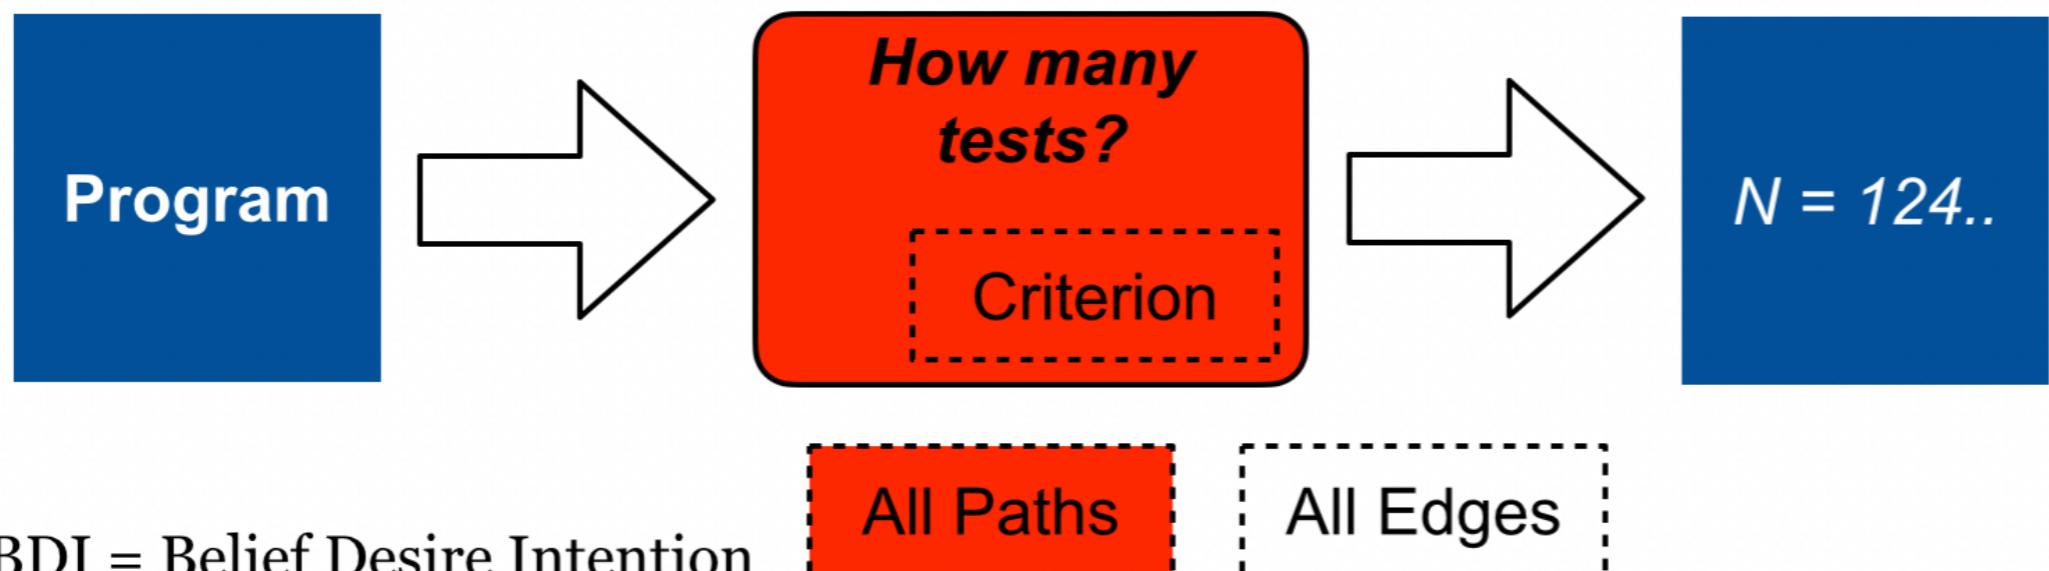
# Assurance ...

- How testable are BDI agents?
- How can I *verify* my MAS?
- How can I *debug* my MAS?
- *Explaining* my agent

# How testable are agents?

Winikoff & Cranefield, 2014

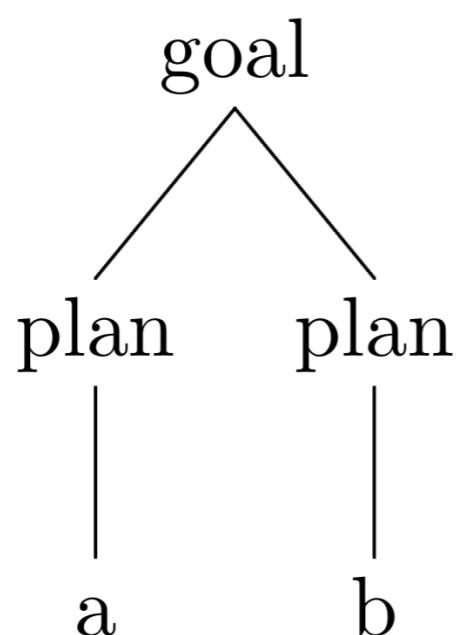
- Assurance important, traditionally done through testing
- But intuitively agents harder to test - is this true?
- Consider all paths testing criterion
- How many paths are there for a BDI\* program?



(\*) BDI = Belief Desire Intention

# BDI Goal-Plan Trees

- Goal: select one of its applicable plan instances
- Plan: execute all sub-goals (in order)
- Failure handling mechanism varies between BDI platforms. Common mechanism: when plan fails, re-post the goal, and select an unused plan instance



Without Failure Handling		With Failure Handling	
a	<b>X</b>	a	<b>X</b>
b	<b>X</b>	b	<b>X</b>
		a <b>X</b> b	a <b>X</b> b <b>X</b>
		b <b>X</b> a	b <b>X</b> a <b>X</b>

Parameters			Number of			No Failure Handling		With Failure Handling	
$j$	$k$	$d$	goals	plans	actions	$n^{\checkmark}(g)$	$n^*(g)$	$n^{\checkmark}(g)$	$n^*(g)$
2	2	3	21	42	62 (13)	128	614	$6.33 \times 10^{12}$	$1.82 \times 10^{13}$
3	3	3	91	273	363 (25)	1,594,323	6,337,425	$1.02 \times 10^{107}$	$2.56 \times 10^{107}$
2	3	4	259	518	776 (79)	1,099,511,627,776	6,523,509,472,174	$1.82 \times 10^{157}$	$7.23 \times 10^{157}$
3	4	3	157	471	627 (41)	10,460,353,203	41,754,963,603	$3.13 \times 10^{184}$	$7.82 \times 10^{184}$

- Astronomical numbers
- Failure handling makes a huge difference

With Failure Handling	
$n^{\checkmark}(g)$	$n^*(g)$
$6.33 \times 10^{12}$	$1.82 \times 10^{13}$
$1.02 \times 10^{107}$	$2.56 \times 10^{107}$
$1.82 \times 10^{157}$	$7.23 \times 10^{157}$
$3.13 \times 10^{184}$	$7.82 \times 10^{184}$

# Reality check

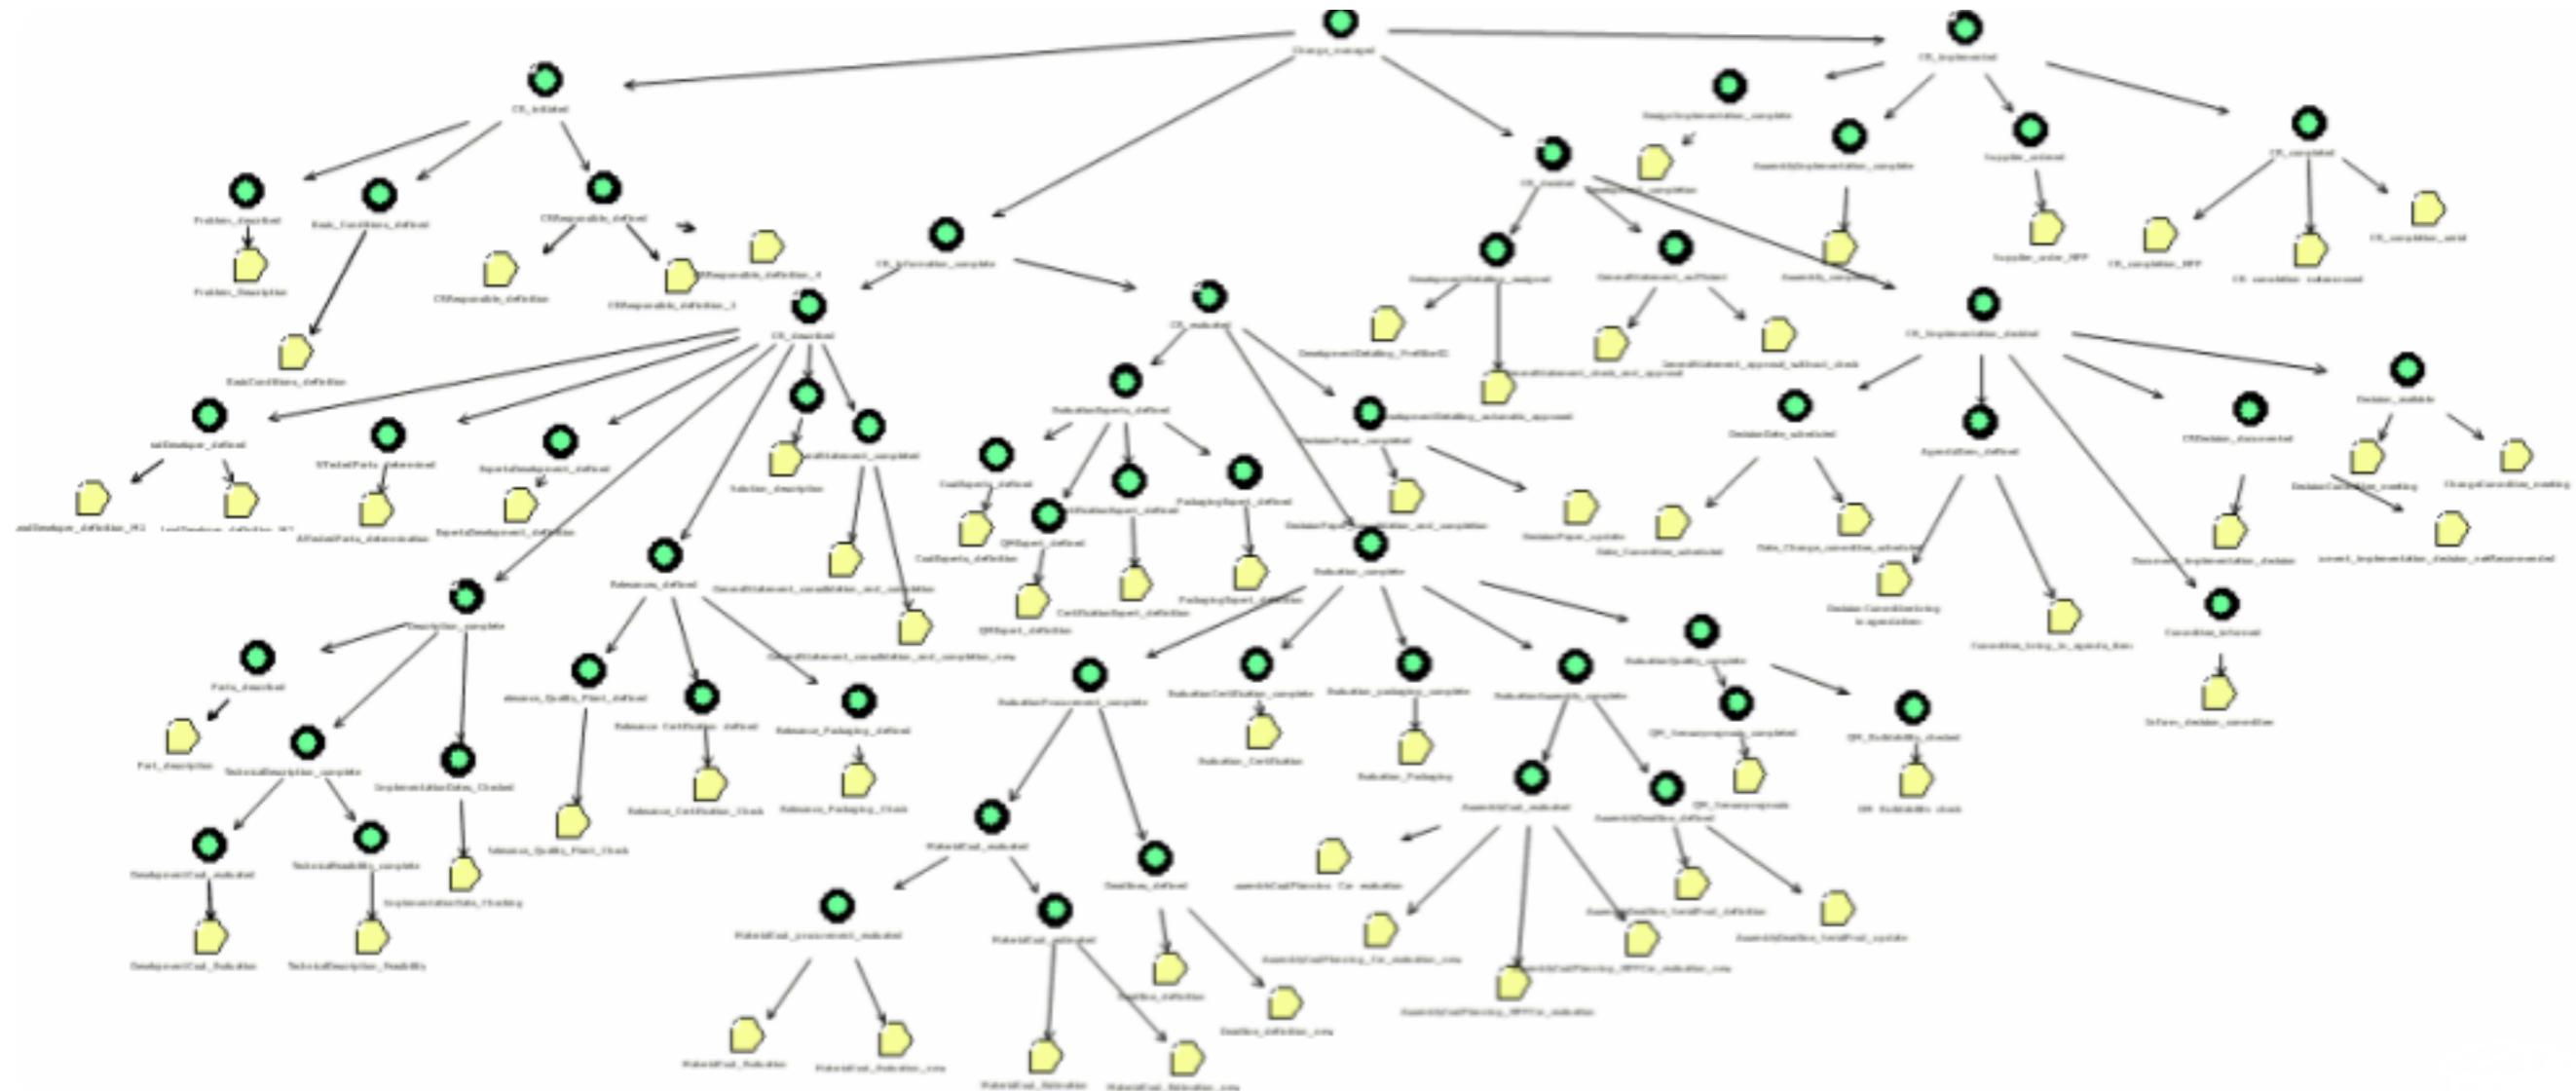


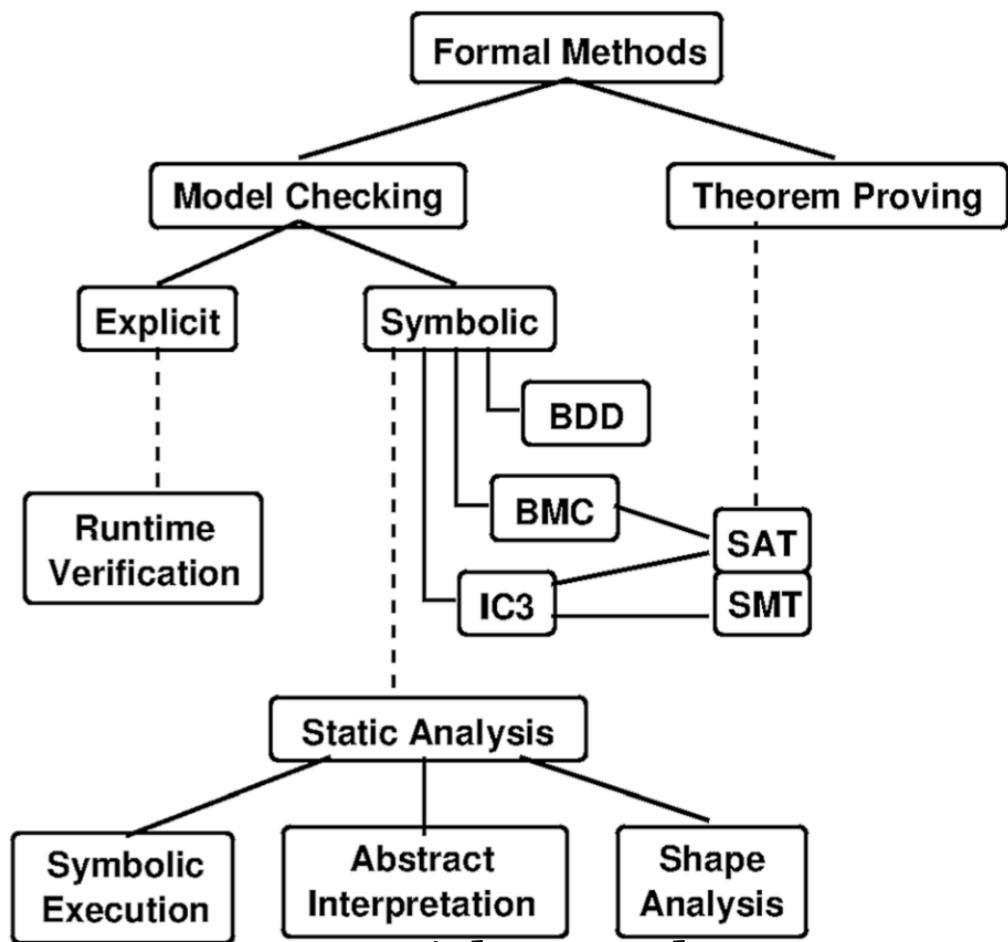
Figure 11: Goal-plan tree from the work of Burmeister et al. (2008, Figure 6) (reproduced with permission from IFAAMAS)

Parameters $j$ $k$ $d$	Number of goals    plans    actions			No Failure Handling		With Failure Handling		$n(m)$
	goals	plans	actions	$n^*(g)$	$n^*(g)$	$n^*(g)$	$n^*(g)$	
2    2    3	21	42	62 (13)	128	614	$6.33 \times 10^{12}$	$1.82 \times 10^{13}$	6,973,568,802
3    3    3	91	273	363 (25)	1,594,323	6,337,425	$1.02 \times 10^{107}$	$2.56 \times 10^{107}$	$5.39 \times 10^{57}$
2    3    4	259	518	776 (79)	1,099,511,627,776	6,523,509,472,174	$1.82 \times 10^{157}$	$7.23 \times 10^{157}$	$2.5 \times 10^{123}$
3    4    3	157	471	627 (41)	10,460,353,203	41,754,963,603	$3.13 \times 10^{184}$	$7.82 \times 10^{184}$	$5.23 \times 10^{99}$
Workflow with 57 goals(*)				294,912	3,250,604	$2.98 \times 10^{20}$	$9.69 \times 10^{20}$	( $\ell = 4$ )
(*) The paper says 60 goals, but their figure 6 actually has 57 goals.				294,912	1,625,302	$6.28 \times 10^{15}$	$8.96 \times 10^{15}$	( $\ell = 2$ )
				294,912	812,651	$9.66 \times 10^{11}$	$6.27 \times 10^{11}$	( $\ell = 1$ )

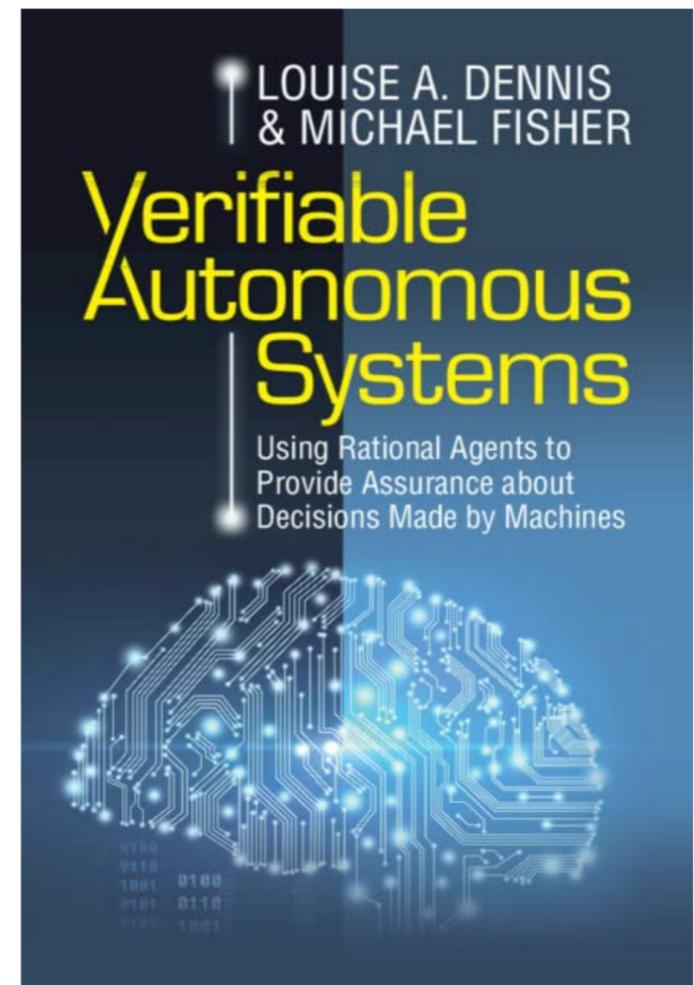
	With Failure Handling			$n(m)$
	$n^*(g)$	$n^*(g)$	$n^*(g)$	
Abstract tree ( $j,k=2-3, d=3$ )	$6.33 \times 10^{12}$	$1.82 \times 10^{13}$	$1.82 \times 10^{157}$	$6.973,568,802$
	$1.02 \times 10^{107}$	$2.56 \times 10^{107}$	$7.23 \times 10^{157}$	$5.39 \times 10^{57}$
	$3.13 \times 10^{184}$	$7.82 \times 10^{184}$	$3.13 \times 10^{184}$	$5.23 \times 10^{99}$
Real industry application	$2.98 \times 10^{20}$	$9.69 \times 10^{20}$	$2.98 \times 10^{20}$	( $\ell = 4$ )
	$6.28 \times 10^{15}$	$8.96 \times 10^{15}$	$6.28 \times 10^{15}$	( $\ell = 2$ )
	$9.66 \times 10^{11}$	$6.27 \times 10^{11}$	$9.66 \times 10^{11}$	( $\ell = 1$ )

# How can I verify my MAS?

- Range of techniques ...
- Formal verification of agent programs (Dennis & Fisher, 2023)
- One key challenge: identifying & capturing requirements! (Next slide)



Source: Fisher et al, 2021



# Initial process

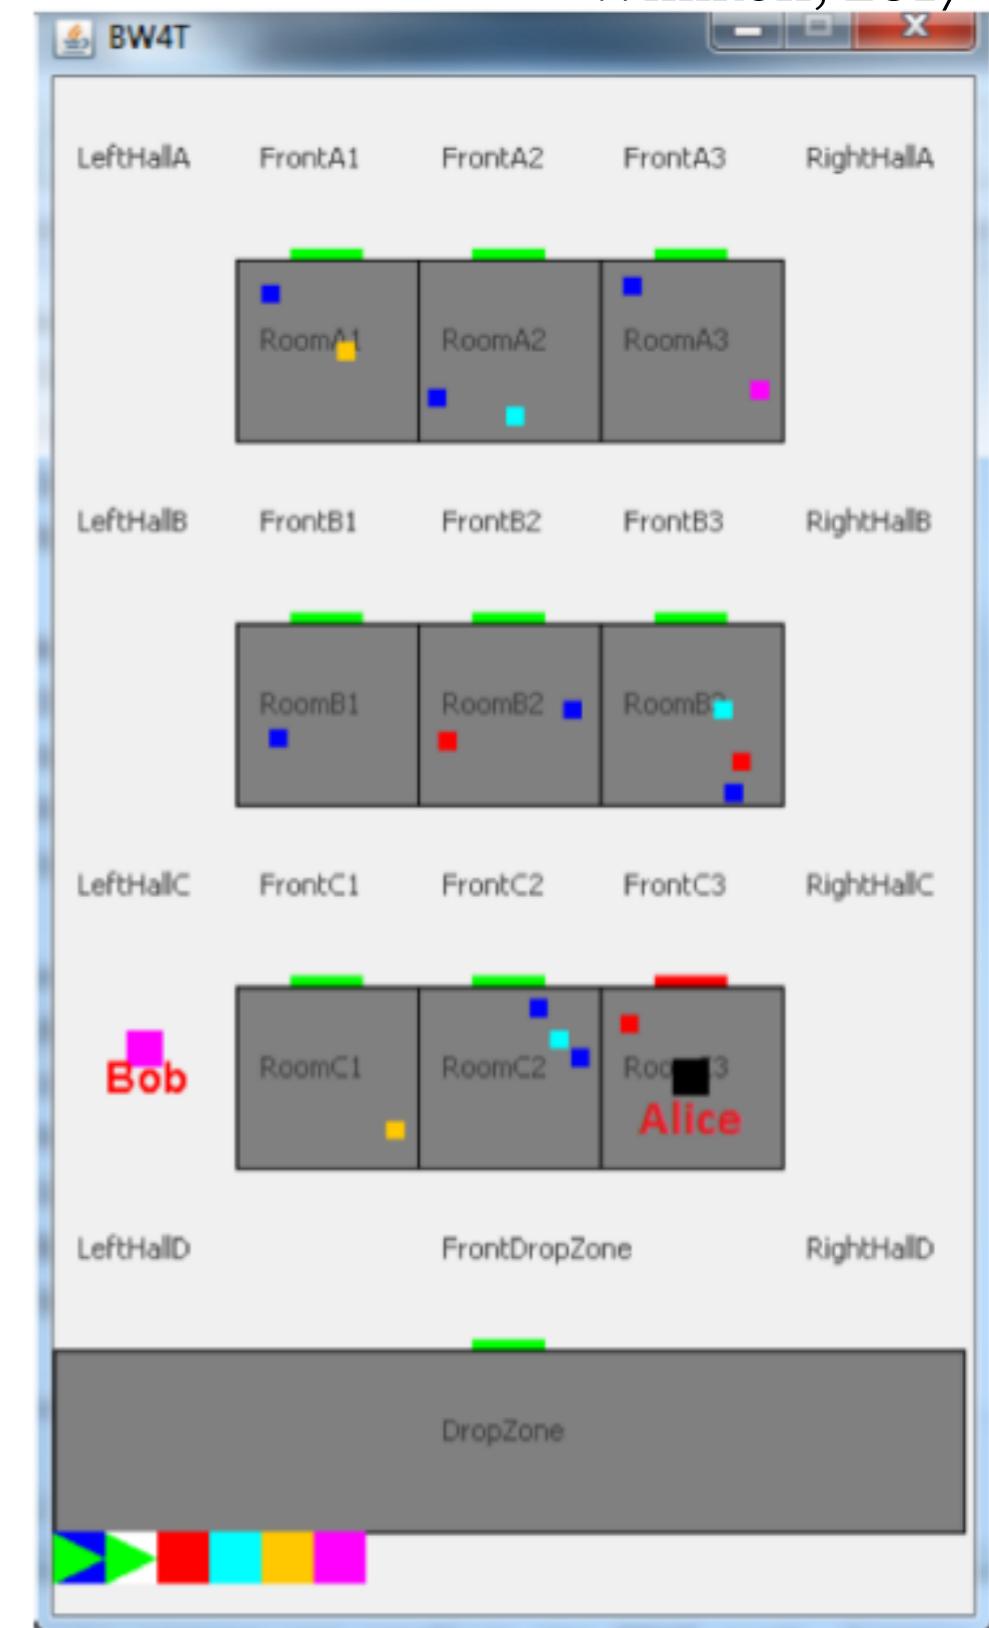
Fisher et al, 2021

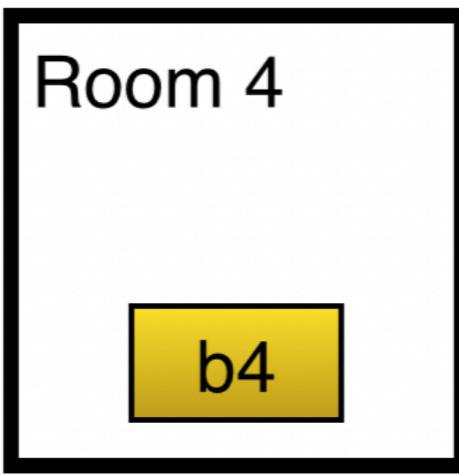
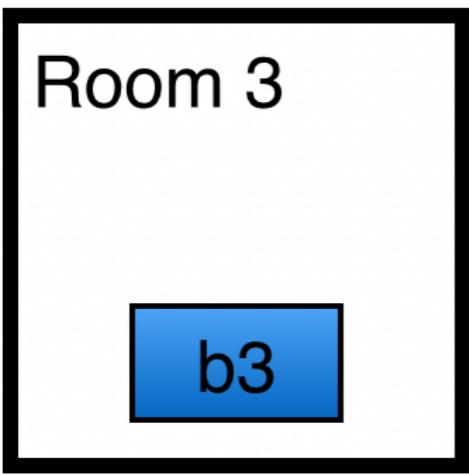
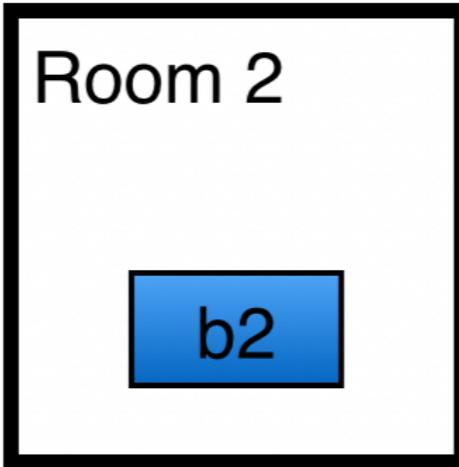
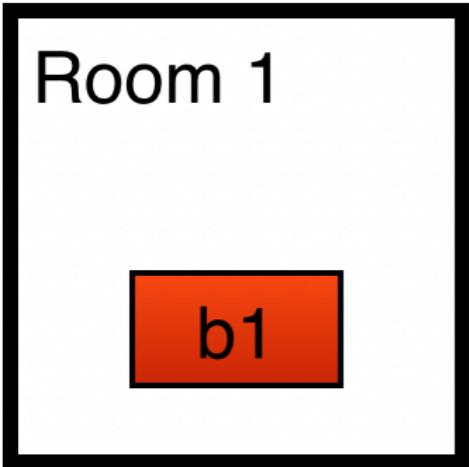
- Consider the **human context** (licensing, assumed human capabilities, laws/regulation, the interface) to **identify requirements**
  - Licensing: physical capabilities, domain knowledge, regulatory knowledge, ethical normalisation
  - Capabilities: general knowledge, common sense, physical capabilities, maturity/life-experience, ethics/values (e.g. self-preservation)
- Consider **how to assess requirements for an autonomous system**; e.g. assess behaviour in scenarios, rather than exam-based assessment of domain knowledge

# How can I *debug* my MAS?

Winikoff, 2017

- Agent programs hard to debug: parallelism, complex control flow
- Key idea: use interactive questioning (“Why?”) (Hindriks 2012, Ko & Myers 2008)
- Example: Blocks World for Teams (BW4T)
- Complementary to algorithmic debugging (Ahlbrecht, 2023)





Goal:    1    2    3

- 1 goto(room1)
- 2 gotoBlock(b1) % red
- 3 pickUp
- 4 goto(dropzone)
- 5 putDown % now need yellow ...
- 6 goto(room2) % exploring ...
- 7 goto(room3)
- 8 goto(room4)
- 9 gotoBlock(b4) % yellow
- 10 pickUp
- 11 goto(dropzone)
- 12 putDown % now need blue ...
- 13 gotoBlock(b3) % blue
- 14 pickUp
- 15 call(deliver)
- 16 gotoBlock(b2) % blue
- 17 pickUp % FAILED (because already holding b3)
- 18 goto(dropzone)
- 19 putDown % now done

+!deliver : nextColour(done)  $\leftarrow$  +done. % If done then stop  
% select a block of the right colour and go get and deliver it

+!deliver : colour(B, C)  $\wedge$  nextColour(C)  $\wedge$   $\neg$ holding(B)  $\leftarrow$  gotoBlock(B) ; pickUp ; !deliver.

+!deliver : holding(B)  $\wedge$  colour(B,C)  $\wedge$  nextColour(C)  $\leftarrow$  goto(dropzone) ; putDown ; !deliver.

% if holding a block that is not the next colour required then put it down (this may occur if e.g. someone else delivers a block, so the next colour changes)

+!deliver : holding(B)  $\wedge$  colour(B,C)  $\wedge$   $\neg$ nextColour(C)  $\leftarrow$  putDown ; !deliver.

% if I know of a place that I've not yet visited then go there (explore)

+!deliver : place(P)  $\wedge$   $\neg$ beenthere(P)  $\leftarrow$  goto(P) ; !deliver.

# Debugging with Questions

Q1: why did pickUp at 17?

A1: because (i) !deliver at 15, (ii) plan  $\pi_2$  has true context condition of "*colour(b2,blue) & nextColour(blue) & ~holding(b2)*", and (iii) the step gotoBlock(b2) succeeded

Q2: why believe colour(b2,blue) & nextColour(blue) & ~holding(b2) at 15?

A2: because putDown done at 12, and colour(b2,blue) & *~holding(b2)* was true

Q3: why believe *~holding(b2)* at 12?

A3: that's been true since the beginning of execution ...

11 goto(dropzone)

**Q3 12 putDown**

13 gotoBlock(b3) % blue  
14 pickUp

**Q2 15 call(deliver)**

16 gotoBlock(b2) % blue

**Q1 17 pickUp % FAILED**

...

$\pi_2 = +\text{!deliver} :$   
 $\text{colour(B, C)}$   
 $\& \text{nextColour(C)}$   
 $\& \text{~holding(B)}$   
 $\leftarrow \text{gotoBlock(B)} ;$   
 $\text{pickUp ; !deliver.}$

# Questions and Answers

- Why (not) do S?
  - Prior step  $S_n$
- Why (not) believe C?
  - Prior belief  $C_n$
- Why did step S succeed/fail?
  - No Applicable Plan
  - “Inevitable”
- Why was plan P selected?
  - “Actually, no”
  - AND, OR

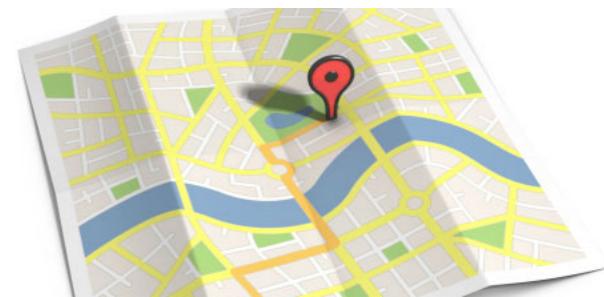
# How to answer “Why?”

*Principle: answer with prior point in execution where a different option would have led to a different outcome.*

- Begin with closest explanation point
- Repeated "Why?" finds earlier explanations ...

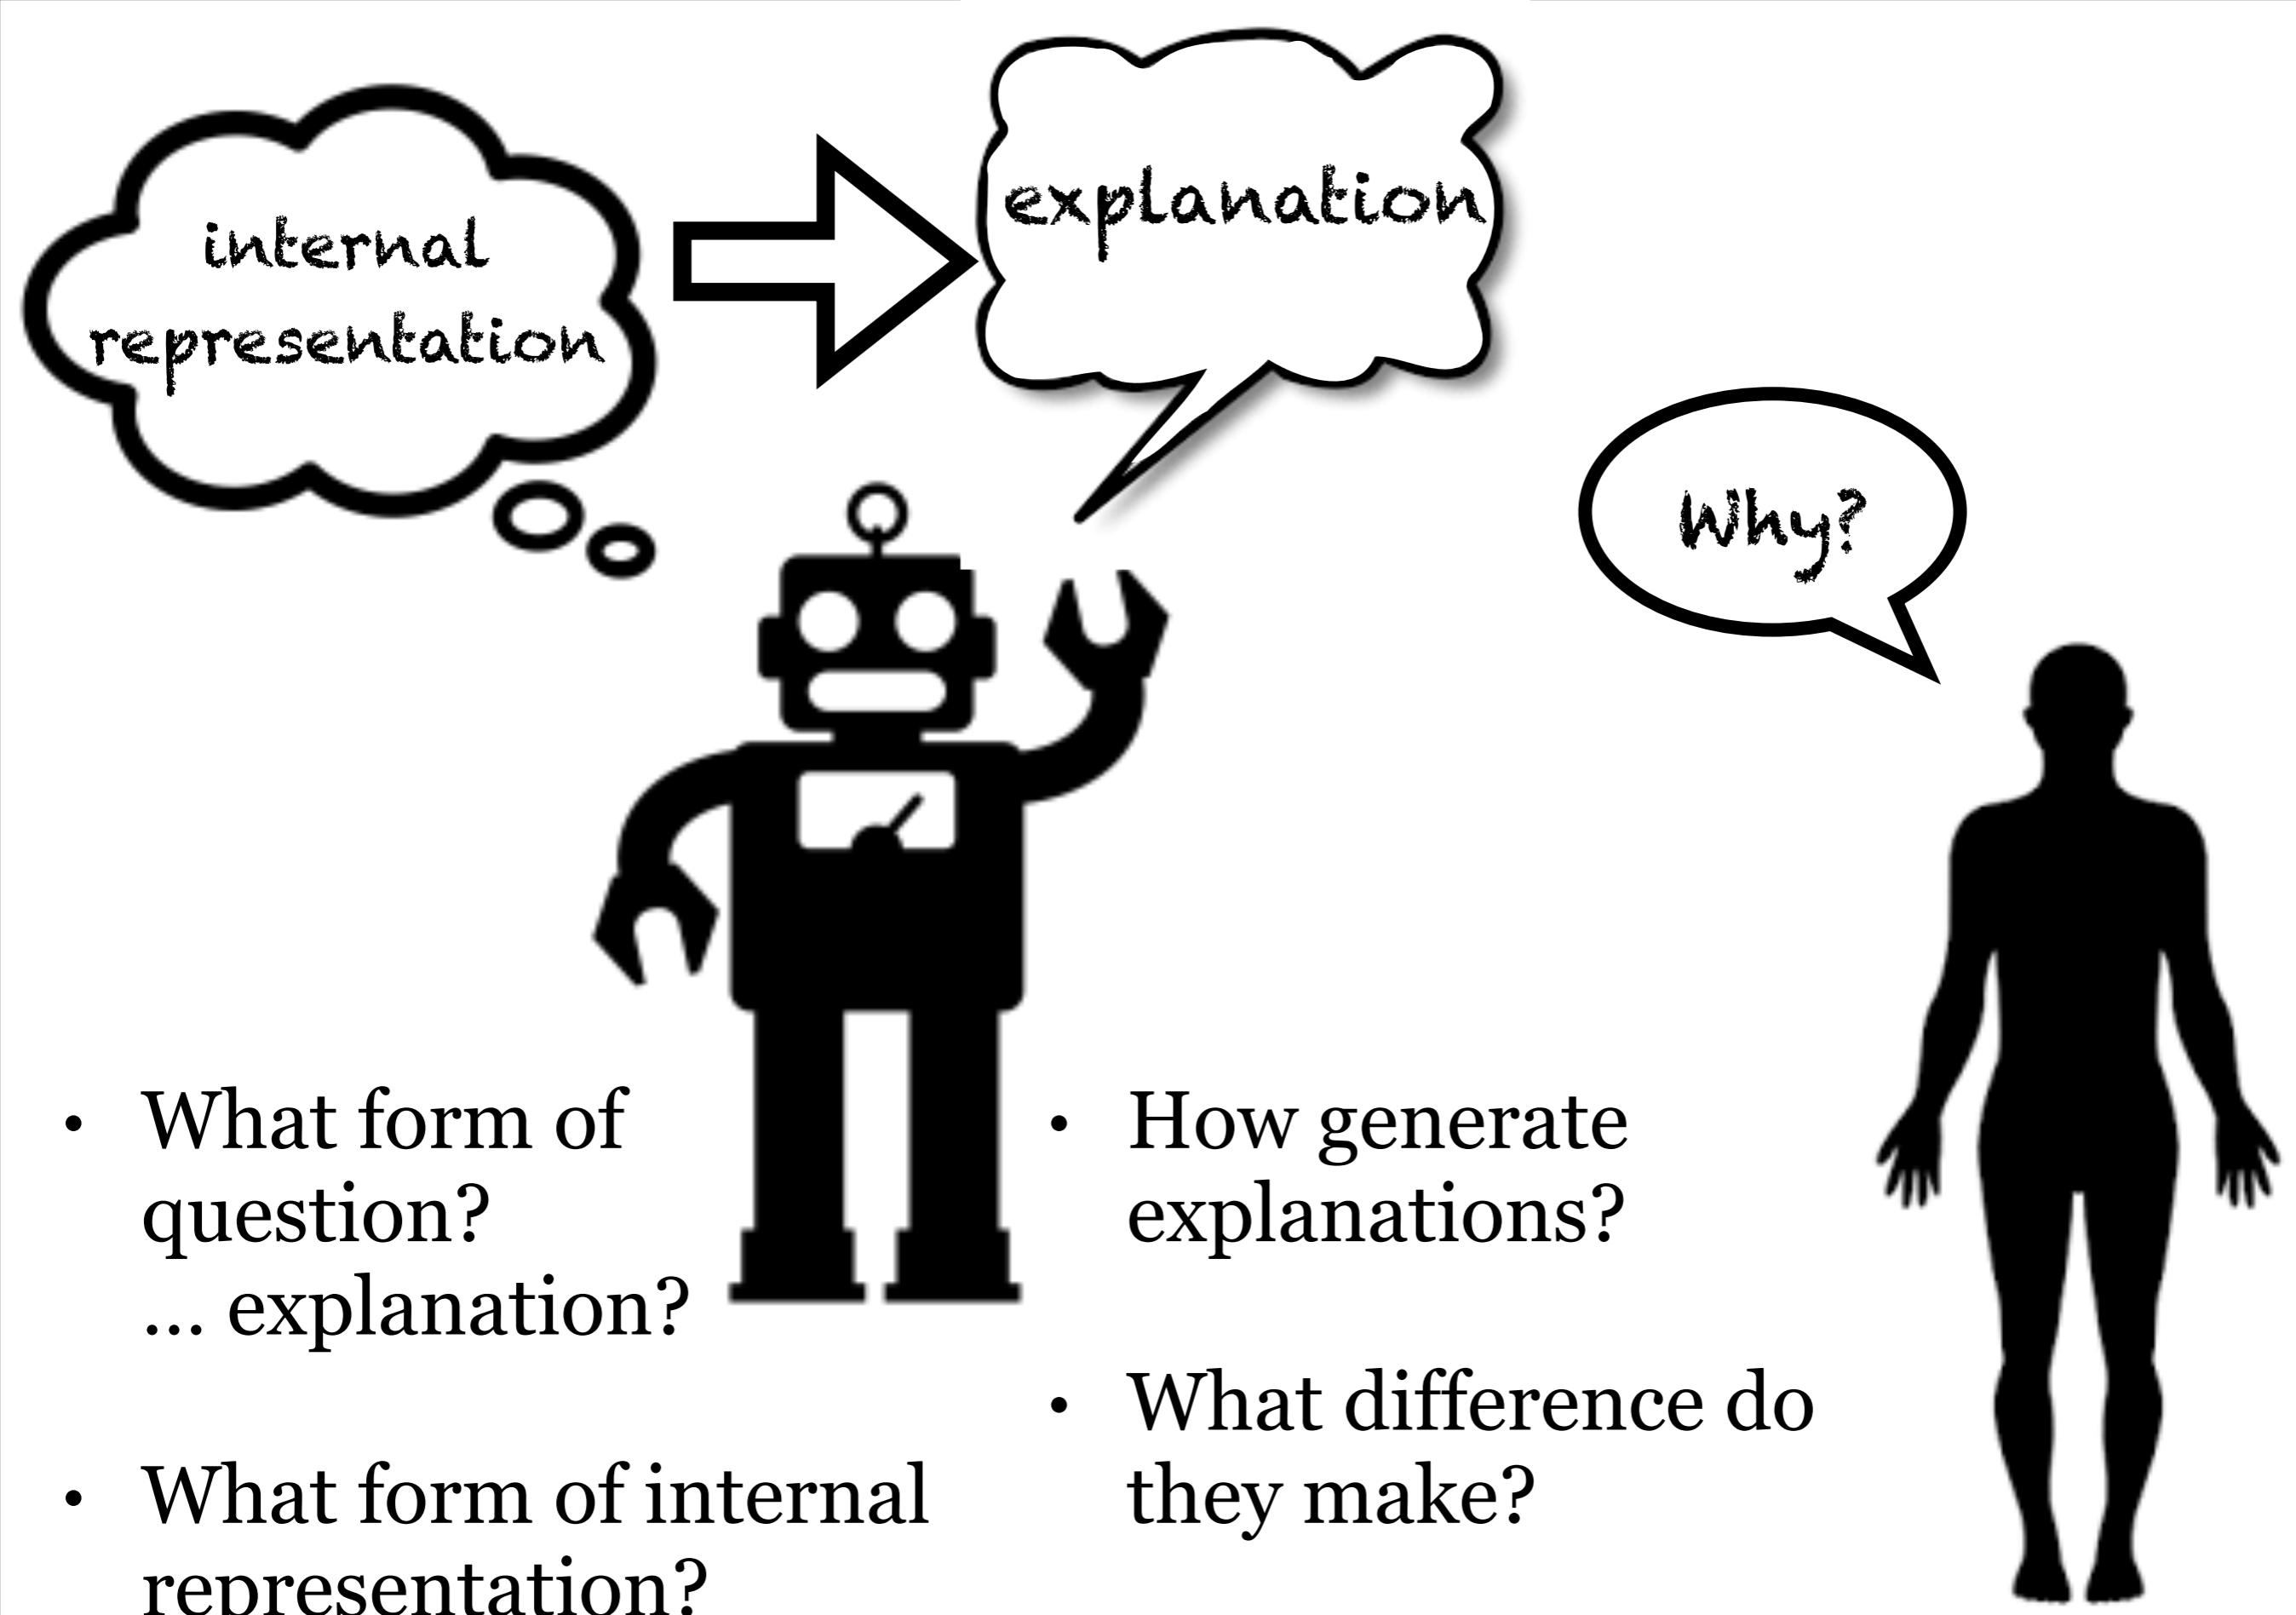
# *Explaining my agent*

- Explain why a decision was made (or course of action taken)
- Important for trust, especially when system finds a non-obvious solution, and for operating in a social world

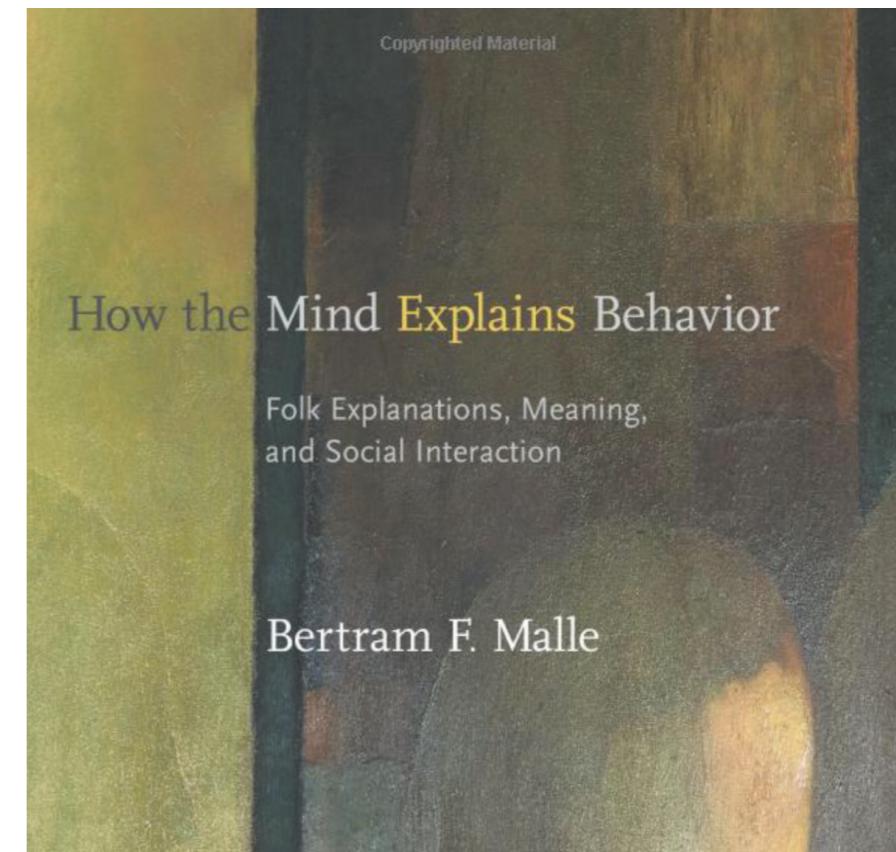


***“An ability of a system to explain its advice was thought to be its most important attribute. Second in importance was the ability of a system to understand and update its own knowledge base. ...” (Teach & Shortliffe, 1981, emphasis added)***

***“... for instance ... for users of care or domestic robots a why-did-you-do-that button which, when pressed, causes the robot to explain the action it just took” (IEEE, Ethically Aligned Design v1, p20)***



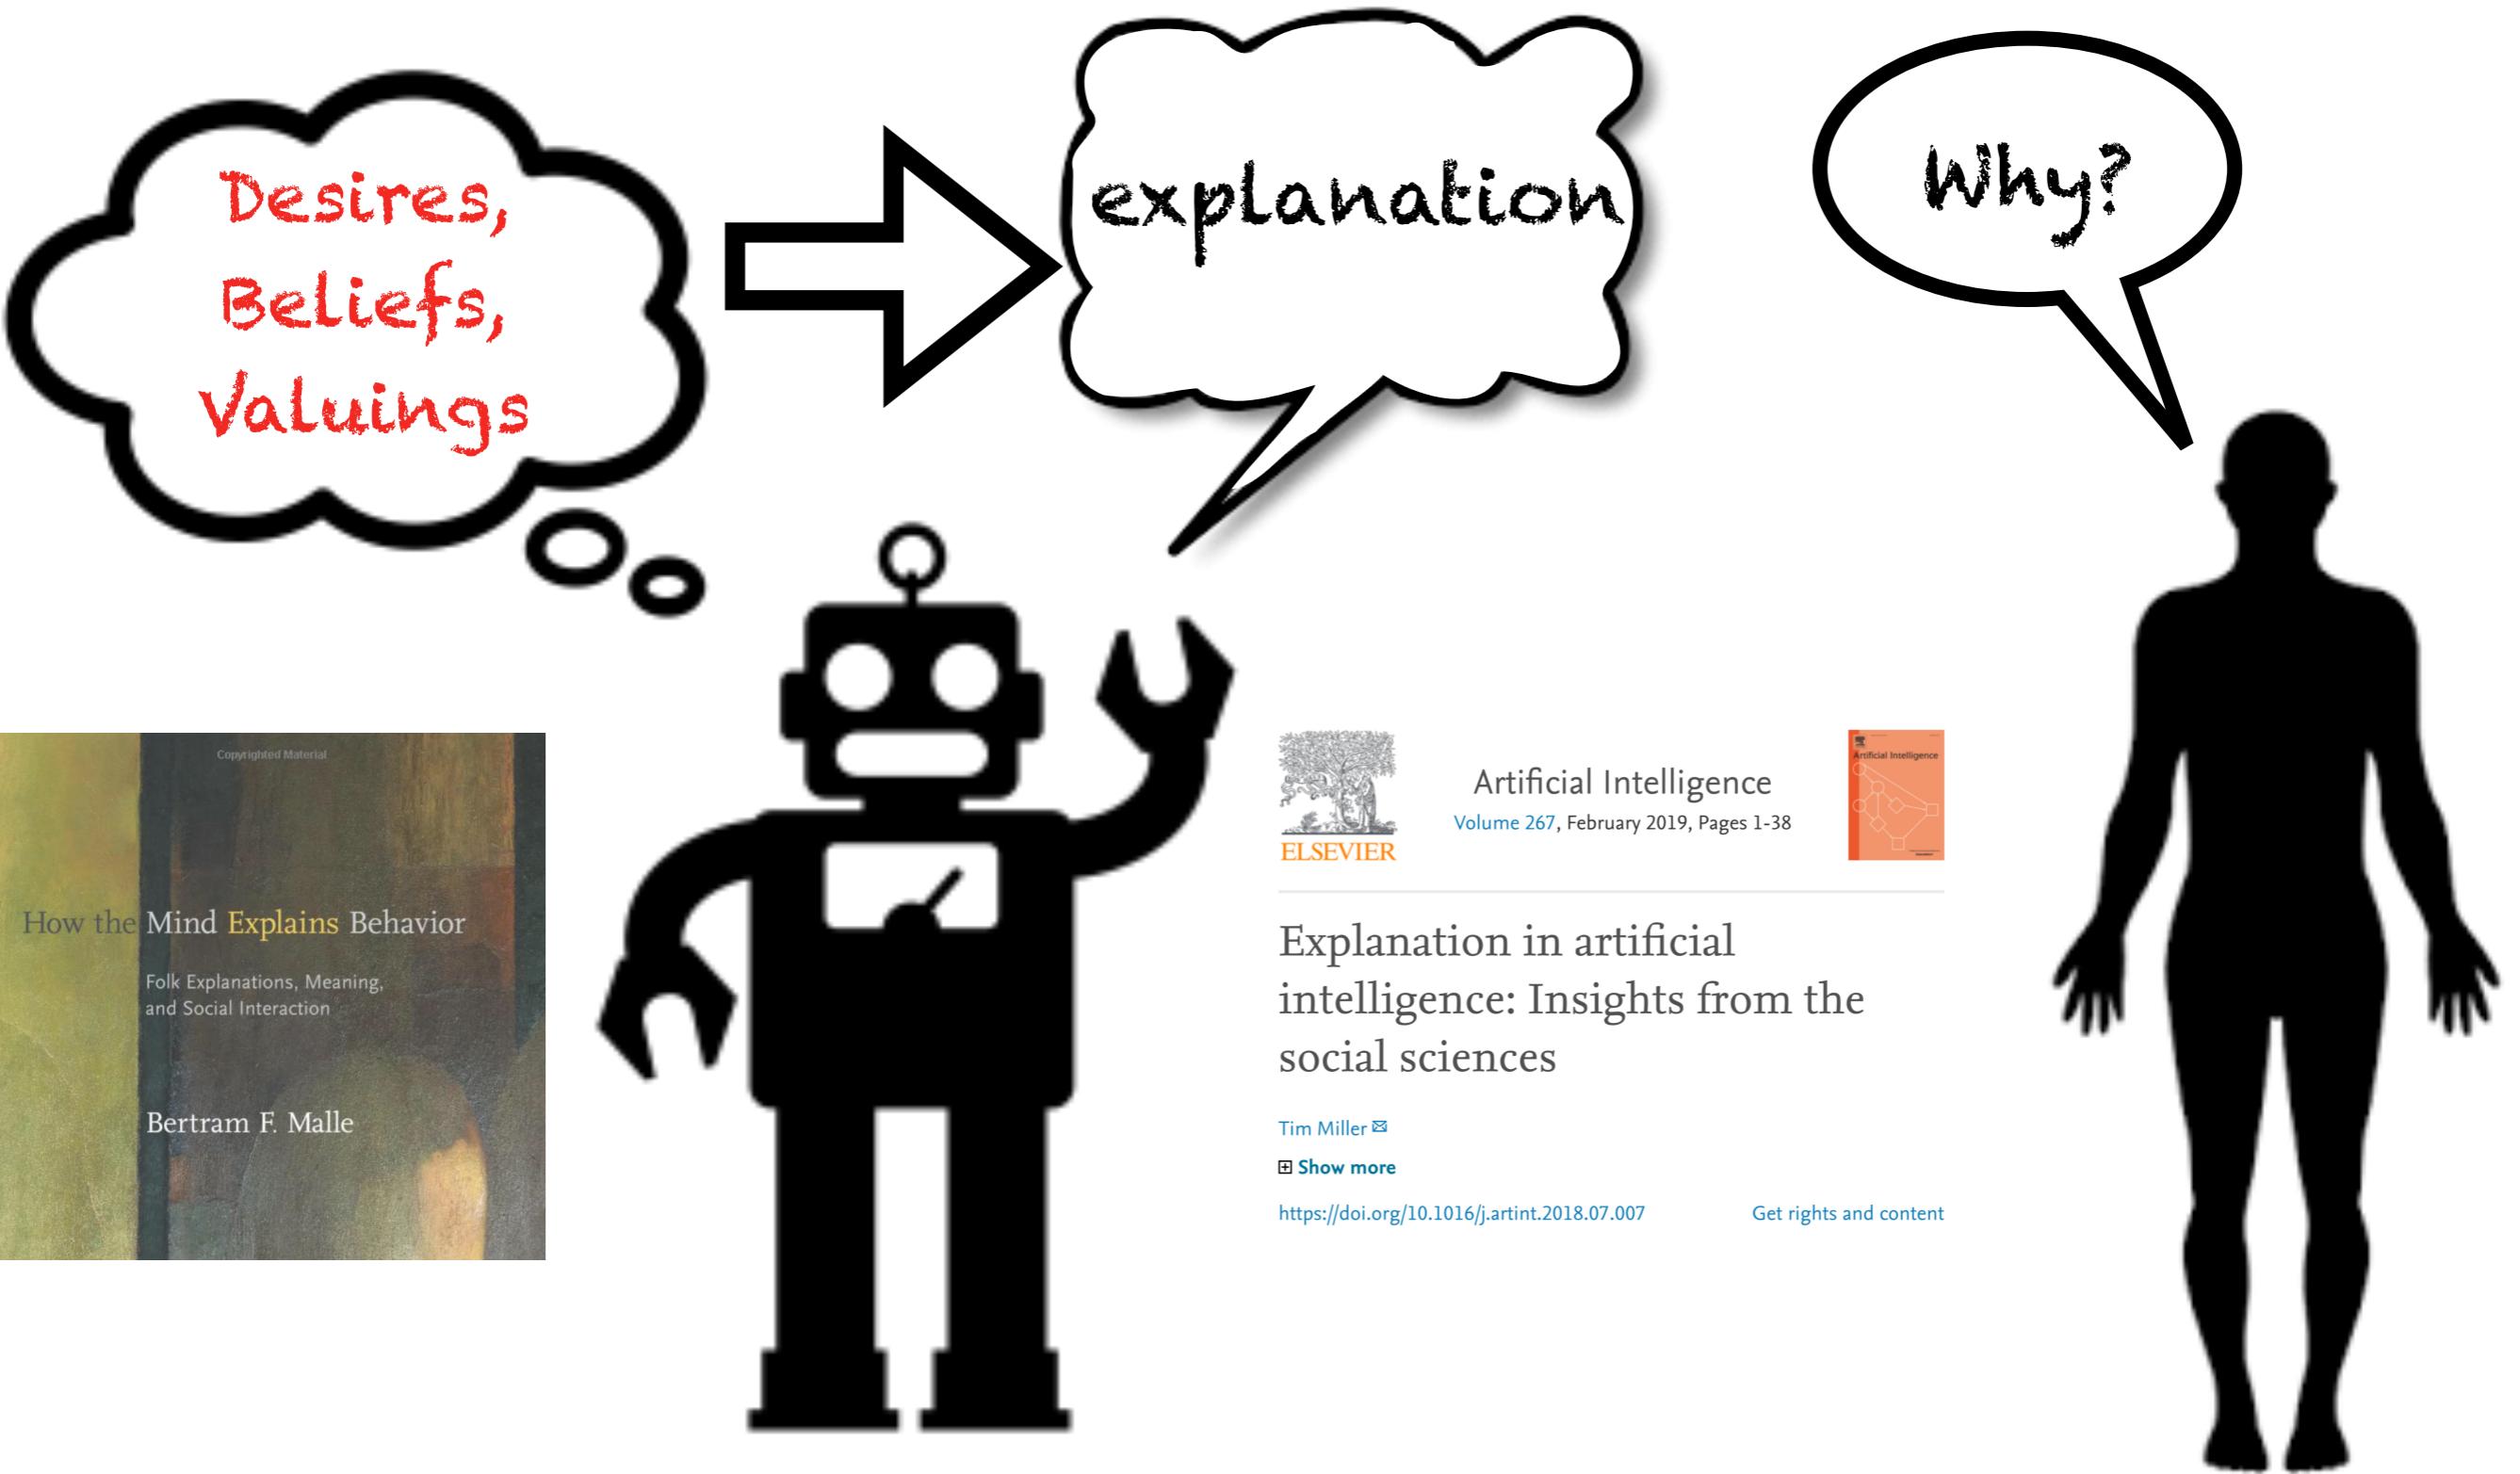
- Human explanations use *folk psychological constructs*, specifically Beliefs, Desires, Valuings (Malle, 2004)
- BDI agents use beliefs, desires ...



# Key claim

Winikoff et al, 2021

Cognitive Agents are **explainable**, since they use similar concepts to those humans use in explanation.



“I had money, Ann was not in her office, I preferred shop coffee to the kitchen (because it’s better quality coffee), I needed to go to the shop in order to do getCoffee(shop), and I desired to getShopCoffee”

# Summary

- EMAS - What? Concepts, notations, processes, techniques, tools, and languages to effectively develop MAS
- EMAS - Why?
- Some challenges: practicalities, motivation, community
- Research: beyond BDI; what about learning?; flexible interactions; assurance (testing, debugging, verification) ... and explanation

# Thank you!

- Abdullah Al-Amin (s)
  - Aditya Ghose (AU)
  - Akın Günay (s\*)
  - Alex Richter (NZ)
  - Alexander Egyed (AT)
  - Alexander Reder (AT)
  - Aloys Mbala (AU)
  - Amir Aryani (s)
  - Andrew Hodgson (AOS, AU)
  - Andrew Lucas (AOS, AU)
  - Angelo Ferrando (s\*, IT)
  - Ayelet Cohen (NZ)
  - Bastin Tony Roy Savarimuthu (NZ)
  - Beatriz López (ES)
  - Bernd-Holger Schlingloff (DE)
  - Brent Martin (NZ)
  - Carles Sierra (ES)
  - Cathal Doyle (NZ)
  - Christoph Matthaei (NZ)
  - Christopher Cheong (s,AU)
  - David Poutakidis (s)
  - David Pym (UK)
  - Dhirendra Singh (AU)
  - Duc Q. Pham (s)
  - Ed Kazmierczak (AU)
  - Frank Dignum (NLSE)
  - Galina Sidorenko (NZSE)
  - Gaya Buddhinath Jayatilleke (s)
  - Guannan (“Peter”) Li (s)
  - Hanno-Felix Wagner (s\*)
  - Harald Søndergaard (AU)
  - Heinz Schmidt (AU)
  - Hywel Lloyd (BPAC, NZ)
  - Ian D. Peake (AU)
  - Ian Mathieson (AU)
  - Islam Elgadawy (s)
  - James Harland (AU)
  - James Harland (AU)
  - Jamie Curmi (AOS, AU)
  - Jane Li (NZ)
  - Jason Khalouf (s)
  - Jenny McDonald (NZ)
  - Jocelyn Cranefield (NZ)
  - John Thangarajah (s,AU)
  - Joshua Hutchison (s)
  - Joy R. Rudland (NZ)
  - Judith Swan (NZ)
  - Julija Sardelić (NZ)
  - June Tordoff (NZ)
  - Khanh Hoa Dam (s,AU)
  - Klaus Fischer (DE)
  - Kristin Yvonne Rozier (US)
  - Lavindra Priyalal de Silva (s,UK)
  - Leon Sterling (AU)
  - Lin Padgham (AU)
  - Louise Dennis (UK)
  - M. Birna van Riemsdijk (NL)
  - Mal Gorman (BoM, AU)
  - Manjula Devananda (s)
  - Margaret Hamilton (AU)
  - Mariusz Nowostawski (NZ NO)
  - Maryam Purvis (NZ)
  - Massimo Cossentino (FR/IT)
  - Mehdi Dastani (NL)
  - Michael Fisher (UK)
  - Minjie Hu (s)
  - Myrthe Tielman (NL)
  - Neil Yorke-Smith (NL)
  - Nitin Yadav (AU)
  - Omer Rana (UK)
  - Paul Maisano (AOS, AU)
  - Peter Vlugter (NZ)
  - Phil Blyth (NZ)
  - Philip Dart (AU)
  - Pinar Yolum (TRNL)
  - Rafael Bordini (UK, BR)
  - Rainer Unland (DE)
  - Ralph Rönnquist (AOS, AU)
  - Richard Zeng (NZ)
  - Rob Wass (NZ)
  - Roberto E. Lopez-Herrejon (AT)
  - Roger Jarquin (Jade, NZ)
  - Sandy Dance (BoM, AU)
  - Sarah Rennie (NZ)
  - Scott DeLoach (US)
  - Sebastian Rodriguez (AU)
  - Sharmila Savarimuthu (NZ)
  - Stephen Cranefield (NZ)
  - Stephen Duffull (NZ)
  - Swee Kin Loke (NZ)
  - Tatjana Lutovac (s\*)
  - Thomas Juan (s\*)
  - Thomas Young (s\*)
  - Tim Miller (AU)
  - Toan Phung (s)
  - Tobias Ahlbrecht (s\*)
  - Virginia Dignum (NLSE)
  - Viviana Mascardi (IT)
  - Wamberto Vasconcelos (UK)
  - Wei Liu (AU)
  - Yoosef Abushark (s\*)
  - Zahir Tari (AU)
  - Zhiyong Zhang (s\*)
- s = student  
s\* = not my student  
XX (e.g. AU) = country  
AOS/BoM/BPAC/JADE = industry

WITH THE RISE OF  
SELF-DRIVING VEHICLES  
THERE WILL EVENTUALLY  
BE A COUNTRY SONG ABOUT  
HOW YOUR TRUCK  
LEFT YOU TOO

<https://youtu.be/ESNDtH6suUo>

<https://youtu.be/mxhcQgvBD2Y>



WITH THE RISE OF  
SELF-DRIVING VEHICLES  
THERE WILL EVENTUALLY  
BE A COUNTRY SONG ABOUT  
HOW YOUR TRUCK  
LEFT YOU TOO

# Bibliography (1/2)

Tobias Ahlbrecht. 2023. An algorithmic debugging approach for belief-desire-intention agents. *Ann Math Artif Intell.* <https://doi.org/ms96>

Roxana A. Belecheanu, Steve Munroe, Michael Luck, Terry Payne, Tim Miller, Peter McBurney, and Michal Pěchouček. 2006. Commercial applications of agents: lessons, experiences and challenges. *AAMAS*. <https://doi.org/c3t6x5>

Steve S. Benfield, Jim Hendrickson, and Daniel Galanti. 2006. Making a strong business case for multiagent technology. *AAMAS*. <https://doi.org/cspvtz>

Rafael H. Bordini, Mehdi Dastani, Jürgen Dix, and Amal El Fallah Seghrouchni. 2005. Multi-Agent Programming: Languages, Platforms and Applications. *Multiagent Systems, Artificial Societies, and Simulated Organizations 15*, Springer, ISBN 0-387-24568-5

Rafael H. Bordini, Mehdi Dastani, Jürgen Dix, and Amal El Fallah Seghrouchni. 2009. Multi-Agent Programming, Languages, Tools and Applications. Springer, ISBN 978-0-387-89298-6

Rafael H. Bordini and Jürgen Dix. 2016. Programming Multiagent Systems, chapter 13 of *Multiagent Systems*, Second Edition Edited by Gerhard Weiss, 2016. ISBN: 9780262533874. MIT Press. Pages 587-639

Rafael H. Bordini, Amal El Fallah Seghrouchni, Koen Hindriks, and Brian Logan, Alessandro Ricci. 2020. Agent programming in the cognitive era. *Auton Agent Multi-Agent Syst.* <https://doi.org/ggwqv2>

Michael E. Bratman. 1987. Intentions, Plans, and Practical Reason. Harvard University Press, Cambridge, MA.

Daniela Briola, Angelo Ferrando, and Viviana Mascardi. 2023. Fantastic MASs and Where to Find Them: First Results and Lesson Learned. *EMAS*. <https://doi.org/ms97>

Amit K. Chopra and Samuel H. Christie V. 2023. Communication Meaning: Foundations and Directions for Systems Research. *AAMAS*.

Louise Dennis and Michael Fisher. 2023. Verifiable Autonomous Systems: Using Rational Agents to Provide Assurance about Decisions Made by Machines. Cambridge University Press, 2023. ISBN 978-1108484992

Virginia Dignum and Frank Dignum. 2010. Designing agent systems: state of the practice. <https://doi.org/fnqs7n>

Michael Fisher, Viviana Mascardi, Kristin Yvonne Rozier, Bernd-Holger Schlingloff, Michael Winikoff, and Neil Yorke-Smith. 2021. Towards a framework for certification of reliable autonomous systems. *Auton Agent Multi-Agent Syst.* <https://doi.org/ms98>

Klaus Havelund, Michael Lowry, SeungJoon Park, Charles Pecheur, John Penix, Willem Visser, and Jonathan White. 2000. Formal Analysis of the Remote Agent Before and After Flight. Proc. 5th NASA Langley Formal Methods Workshop, Williamsburg, VA, Jun. 13-15 2000.

Koen V. Hindriks. 2012. Debugging Is Explaining. *PRIMA* <https://doi.org/ms99>

Koen V. Hindriks. 2014. The Shaping of the Agent-Oriented Mindset. *EMAS*. <https://doi.org/mtbb> Slides: <https://emas14.wordpress.com/wp-content/uploads/2013/11/koen-emas-14.pdf>

The IEEE Global Initiative on Ethics of Autonomous and Intelligent Systems. Ethically Aligned Design: A Vision for Prioritizing Human Well-being with Autonomous and Intelligent Systems, Version 2. IEEE, 2017. [http://standards.ieee.org/develop/indconn/ec/autonomous\\_systems.html](http://standards.ieee.org/develop/indconn/ec/autonomous_systems.html)

F. F. Ingrand, M. P. Georgeff and A. S. Rao. 1993. An architecture for real-time reasoning and system control. *IEEE Expert*. <http://doi.org/dz574s>

Amy J. Ko and Brad A. Myers. 2008. Debugging reinvented: asking and answering why and why not questions about program behavior. *ICSE*. <https://doi.org/dgvqm8>

Brian Logan. 2015. A Future for Agent Programming. *EMAS*. <https://doi.org/mtbc>

Bertram F. Malle. 2004. How the Mind Explains Behavior: Folk Explanations, Meaning, and Social Interaction. The MIT Press. <https://doi.org/mtbd>

# Bibliography (2/2)

Aditya Mathur. 2008. Foundations of Software Testing. Pearson.

Tim Miller. 2019. Explanation in artificial intelligence: Insights from the social sciences. AIJ. <https://doi.org/gfwcxw>

Jörg P. Müller and Klaus Fischer. 2014. Application Impact of Multi-agent Systems and Technologies: A Survey. In: Shehory, O., Sturm, A. (eds) Agent-Oriented Software Engineering. Springer, Berlin, Heidelberg. <https://doi.org/mtbg>

Steve Munroe, Tim Miller, Roxana Belecheanu, Michal Pěchouček, Peter McBurney, and Michael Luck. 2017. Crossing the agent technology chasm: Lessons, experiences and challenges in commercial applications of agents. The Knowledge Engineering Review. <http://doi.org/b6xbv3>

Lin Padgham, John Thangarajah, and Michael Winikoff. 2008. The Prometheus Design Tool – A Conference Management System Case Study. AOSE. <https://doi.org/fn7r5q>

Lin Padgham and Michael Winikoff. Developing Intelligent Agent Systems: A Practical Guide. June 2004, ISBN 0-470-86120-7, John Wiley and Sons.

Alessandro Ricci. 2022. “Go to the Children”: Rethinking Intelligent Agent Design and Programming in a Developmental Learning Perspective. AAMAS.

Sebastian Rodriguez, John Thangarajah, and Michael Winikoff. 2023. A Behaviour-Driven Approach for Testing Requirements via User and System Stories in Agent Systems. AAMAS.

Dhirendra Singh and Lin Padgham. 2015. Community Evacuation Planning for Bushfires Using Agent-Based Simulation: Demonstration. AAMAS.

Munindar P. Singh. 2011. Information-driven interaction-oriented programming: BSPL, the blindingly simple protocol language. AAMAS.

R. L. Teach and E. H. Shortliffe. 1981. An Analysis of Physician Attitudes Regarding Computer-Based Clinical Consultation Systems. Computers and Biomedical Research, 14:542–558.

John Thangarajah, Lin Padgham, and Michael Winikoff. 2005. Prometheus design tool. AAMAS. <https://doi.org/bhchsb>

Michael Winikoff. 2017. Debugging Agent Programs with Why? Questions. AAMAS.

Michael Winikoff and Stephen Cranefield. 2014. On the testability of BDI agent systems. JAIR. <http://doi.org/mtbj>

Michael Winikoff, Lin Padgham, and James Harland. 2001. Simplifying the Development of Intelligent Agents. AI. <https://doi.org/dvbqgb>

Michael Winikoff, Galina Sidorenko, Virginia Dignum, and Frank Dignum. 2021. Why bad coffee? Explaining BDI agent behaviour with valuations. AIJ. <https://doi.org/gm6wss>

Michael Winikoff, Nitin Yadav, and Lin Padgham. 2018. A new Hierarchical Agent Protocol Notation. Auton Agent Multi-Agent Syst. <https://doi.org/gct7gn>

Pınar Yolum and Munindar P. Singh. 2002. Commitment Machines. ATAL. <https://doi.org/c8d6cx>